

---

# Low-Rank Orthogonalization for Large-Scale Matrix Optimization with Applications to Foundation Model Training

Chuan He

Department of Mathematics  
Linköping University, Sweden

chuan.he@liu.se

Zhanwang Deng\*

Academy for Advanced Interdisciplinary Studies  
Peking University, Beijing, People's Republic of China

dzw\_opt2022@stu.pku.edu.cn

Zhaosong Lu

Department of Industrial and Systems Engineering  
University of Minnesota, USA

zhaosong@umn.edu

## Abstract

Neural network (NN) training is inherently a large-scale matrix optimization problem, yet the matrix structure of NN parameters has long been overlooked. Recently, the optimizer Muon (Jordan et al.), which explicitly exploits this structure, has gained significant attention for its strong performance in foundation model training. A key component contributing to Muon's success is matrix orthogonalization. In this paper, we propose *low-rank orthogonalization*, which performs orthogonalization by leveraging the low-rank nature of gradients during NN training. Building on this, we introduce low-rank matrix-signed gradient descent (MSGD) and a low-rank variant of Muon. Numerical experiments show that this low-rank approach can reduce the cost of the matrix-sign step and that low-rank Muon is competitive with, and often improves upon, vanilla Muon in our GPT-2 and LLaMA pretraining experiments. Theoretically, we establish the iteration complexity of low-rank MSGD for finding an approximate stationary solution and, for a residual-controlled low-rank Muon variant used for the theoretical analysis, the iteration complexity for finding an approximate stochastic stationary solution under heavy-tailed noise. The code of this paper can be found in <https://anonymous.4open.science/r/Low-rank-Muon-75C6/..>

**Keywords:** Orthogonalization, Muon, foundation model training, iteration complexity, heavy-tailed noise

## 1 Introduction

Training neural networks (NNs) (LeCun et al., 2015), particularly recent foundation models, has consistently posed challenging large-scale optimization problems. Over the past decade, NN training has been dominated by vector-variate optimization methods—including SGD (Bottou, 2010), AdaGrad (Duchi et al., 2011), RMSprop (Hinton et al., 2012), Adadelata (Zeiler, 2012), Adam (Kingma & Ba, 2015), and AdamW (Loshchilov & Hutter, 2019). Nonetheless, these methods disregard the inherent matrix structure of NN parameters—such as those in multi-layer perceptrons (Rosenblatt, 1958), convolutional layers (LeCun et al., 1989), and the query, key, and value projections in attention mechanisms (Vaswani et al., 2017).

Recently, a shift has occurred: optimization methods that exploit matrix structure are receiving increasing attention and have begun to demonstrate strong performance in foundation model training (Jordan et al.; Pethick et al., 2025). These methods focus on solving the matrix optimization problem:

$$\min_{X \in \mathbb{R}^{m \times n}} f(X). \quad (1)$$

---

\*The first two authors contributed equally.

In particular, Shampoo, as developed in (Anil et al., 2020; Gupta et al., 2018), applies left and right preconditioning matrices and updates them as follows:

$$X^{k+1} = X^k - \eta_k (L^k)^{-1/4} G^k (R^k)^{-1/4}, \quad (2)$$

where  $\eta_k > 0$  is the step size,  $G^k \in \mathbb{R}^{m \times n}$  denotes the (stochastic) gradient of  $f$  at  $X^k$ , and  $L^k \in \mathbb{R}^{m \times m}$  and  $R^k \in \mathbb{R}^{n \times n}$  are the left and right preconditioning matrices, respectively. Shampoo updates the left and right preconditioners  $\{L^k\}$  and  $\{R^k\}$  using the second-order statistics of the accumulated gradients, similarly to AdaGrad, and has shown comparable performance to popular vector-variate optimizers such as Adam and AdamW on foundation model training. Following Shampoo, other matrix-variate optimizers that use two-sided preconditioners, such as CASPR (Duvvuri et al., 2024) and SOAP (Vyas et al., 2025), were also developed. A preconditioned Riemannian gradient descent method was developed in (Bian et al., 2024) for low-rank matrix recovery, which adopts only the diagonal part of the Shampoo preconditioners. Moreover, one-sided preconditioned variants of Shampoo were developed in (An et al., 2025; Xie et al., 2025).

In addition to Shampoo and its variants, another matrix-variate optimizer, Muon (Jordan et al.), has attracted significant attention for outperforming standard optimizers such as Adam and AdamW in foundation model training (AI et al., 2025). At each iteration, Muon performs the update:

$$M^k = (1 - \theta_{k-1}) M^{k-1} + \theta_{k-1} G(X^k; \xi^k), \quad X^{k+1} = X^k - \eta_k \text{msgn}(M^k), \quad (3)$$

where  $G(\cdot; \xi)$  denotes the stochastic gradient of  $f(\cdot)$ , and  $\text{msgn}(M^k) = U^k (V^k)^T$  denotes the matrix sign of  $M^k$ , with  $U^k$  and  $V^k$  containing the left and right singular vectors of  $M^k$ , respectively. The matrix sign computation is often referred to as matrix orthogonalization, because calculating  $\text{msgn}(M)$  is equivalent to finding the (semi-)orthogonal matrix closest to  $M$  with respect to the Frobenius norm (see, e.g., (Bernstein & Newhouse, 2024, Proposition 4)). Muon’s empirical success has sparked significant research interest, including efforts to understand its relationship with other algorithms, establish its convergence guarantees, and propose new variants (see, e.g., (An et al., 2025; Cesista, 2025; Chen et al., 2025; Glentis et al., 2025; Kovalev, 2025; Lau et al., 2025; Li & Hong, 2025; Liu et al., 2025; Ma et al., 2024; Pethick et al., 2025; Riabinin et al., 2025; Sato et al., 2025; Sfyraiki & Wang, 2025; Shen et al., 2025)). A popular interpretation of Muon is from the perspective of a linear minimization oracle with respect to the spectral norm (e.g., see (Bernstein & Newhouse, 2024; Cesista, 2025; Glentis et al., 2025; Kovalev, 2025; Lau et al., 2025; Riabinin et al., 2025)). That is, the matrix sign computation in (3) can be recast as  $-\text{msgn}(M^k) = \arg \min_{\|\Delta\|_{\Delta} \leq 1} \{\langle M^k, \Delta \rangle\}$ , where  $\|\cdot\|$  denotes spectral norm. Based on this interpretation, algorithmic designs leveraging general matrix-induced norms  $\|\cdot\|_{p \rightarrow q}$  have been discussed in (Bernstein & Newhouse, 2024; Cesista, 2025; Glentis et al., 2025; Riabinin et al., 2025). In addition, Muon is also connected to earlier algorithms and can be viewed as a special case of Shampoo, despite not explicitly using preconditioners. As discussed in (Jordan et al.), by taking  $L^k = G^k (G^k)^T$  and  $R^k = (G^k)^T G^k$  in (2), the Shampoo updates reduce to the matrix-signed update:  $X^{k+1} = X^k - \eta_k \text{msgn}(G^k)$ . Furthermore, convergence guarantees for Muon have been extensively studied (e.g., see (An et al., 2025; Chen et al., 2025; Kovalev, 2025; Li & Hong, 2025; Riabinin et al., 2025; Sato et al., 2025; Sfyraiki & Wang, 2025; Shen et al., 2025)), and numerous new variants—such as SWAN (Ma et al., 2024), Scion (Pethick et al., 2025), Gluon (Riabinin et al., 2025), PolarGrad (Lau et al., 2025), Dion (Ahn & Xu, 2025; Ahn et al., 2025), and AdaMuon (Si et al., 2025)—have been proposed. Our analysis complements these recent Muon convergence results rather than replacing them: existing Muon analyses typically study exact or full-rank matrix-sign updates, whereas our theory focuses on inexact low-rank orthogonalization with explicit approximation errors. In the limiting case where the low-rank approximation error vanishes, our framework recovers the full-rank matrix-sign update and is therefore consistent with the existing Muon literature.

Beyond applying orthogonalization, a key innovation of Muon is its use of a GPU-friendly method—Newton-Schulz iterations (typically with five steps)—to perform inexact matrix orthogonalization, making it well-suited for modern foundation model training. In fact, several earlier methods, including spectral gradient descent (Carlson et al., 2015a;b;c) and orthogonalized gradient descent (Tuddenham et al., 2022), have already adopted SVD-based orthogonalization for matrix optimization problems. However, since SVD is not computationally efficient on GPUs when training large-scale neural networks, these methods fail to scale to foundation model training.

---

Inspired by Muon and its orthogonalization subroutine, we aim to develop a faster and more lightweight orthogonalization method to further enhance Muon and its variants. Specifically, our design leverages the widely observed phenomenon that the gradient matrices of NN parameters are often low-rank (see, e.g., (Hao et al., 2024; Hu et al., 2022; Malladi et al., 2023; Zhao et al., 2024); for LoRA–Muon finetuning, see (Bogachev et al., 2026)). To exploit this low-rank property, we propose performing low-rank orthogonalization by incorporating well-known low-rank matrix approximation techniques (Drineas et al., 2006a; Halko et al., 2011). Our approach first constructs a low-rank projection of the gradient matrix using QR decomposition on a sketched matrix, and then performs orthogonalization on the projected matrix by leveraging its structure. This is different from low-rank adaptation methods such as LoRA, which impose low-rank structure on the parameter update itself; here, the low-rank structure is used only inside the optimizer’s orthogonalization subroutine. Our proposed low-rank orthogonalization offers two main advantages over existing orthogonalization methods:

1. **Computational Efficiency:** Traditional orthogonalization can be seen as computing the polar factor of a given full matrix. To exploit the low-rank property, our low-rank orthogonalization first computes the  $Q$  factor of a smaller sketched matrix, and then computes the polar factor of a projected matrix constructed using the  $Q$  factor. Since both the QR decomposition and polar decomposition are performed on much smaller matrices, our low-rank approach enjoys substantial computational savings for large-scale problems.
2. **Noise Robustness:** In the presence of noise, singular vectors associated with small singular values often vary significantly, leading to instability when directly applying orthogonalization to the full matrix. To circumvent this instability, our low-rank orthogonalization method clips these vectors to eliminate unreliable estimates, thereby stabilizing the orthogonalization process and yielding a robust estimate of the matrix sign.

These advantages will be illustrated in detail in Sections 3.1 and 4. Based on low-rank orthogonalization, we develop low-rank matrix-signed gradient descent (MSGD) and a low-rank variant of Muon. We also establish their complexity guarantees under mild assumptions.

Our main contributions are highlighted below.

- We propose *low-rank orthogonalization* that can be readily incorporated into matrix-variate optimization algorithms such as Muon. It can be efficiently executed on GPUs and serves as a lightweight substitute for traditional orthogonalization methods that are directly applied to the full matrix, such as Newton-Schulz iterations.
- Under mild assumptions, we establish the iteration complexity of low-rank MSGD and of a residual-controlled low-rank Muon variant used for the theoretical analysis. To the best of our knowledge, our results provide the first complexity guarantees for a broad class of inexact Muon-type algorithms under heavy-tailed noise.

The remainder of this paper is organized as follows. In Section 2, we introduce the notation and assumptions used throughout the paper. In Section 3, we propose low-rank orthogonalization and, based on it, develop low-rank MSGD and low-rank Muon. In Section 4, we present numerical results. Finally, we provide the proofs of the main results in Section 6.

## 2 Notation and Assumptions

Throughout this paper, we use  $\mathbb{R}^{m \times n}$  to denote the Euclidean space of  $m \times n$  real matrices, and  $\mathbb{Z}_+$  to denote the set of all nonnegative integers. We use  $\|\cdot\|$  to denote the Euclidean norm of a vector or the spectral norm of a matrix;  $\|\cdot\|_*$  and  $\|\cdot\|_F$  to denote the nuclear norm and the Frobenius norm of a matrix, respectively; and  $\langle \cdot, \cdot \rangle$  to denote the trace inner product for matrices. For any  $M \in \mathbb{R}^{m \times n}$ , we use  $\text{rank}(M)$  to denote its rank, and  $[M]_k$  to denote its best rank- $k$  approximation with respect to  $\|\cdot\|_F$ . We define the matrix sign of any nonzero matrix  $M \in \mathbb{R}^{m \times n}$  as  $\text{msgn}(M) = UV^T$ , where  $U \in \mathbb{R}^{m \times r}$  and  $V \in \mathbb{R}^{n \times r}$  are column-orthogonal

matrices obtained from the reduced SVD of  $M$ . We let  $\varrho := \min\{m, n\}$ . In addition, we use  $\tilde{\mathcal{O}}(\cdot)$  to denote  $\mathcal{O}(\cdot)$  with logarithmic factors omitted.

We now make the following assumption throughout this paper.

**Assumption 1.** (a) *There exists a finite  $f_{\text{low}}$  such that  $f(X) \geq f_{\text{low}}$  for all  $X \in \mathbb{R}^{m \times n}$ .*

(b) *There exists an  $L_* > 0$  such that  $\|\nabla f(X) - \nabla f(Y)\|_* \leq L_* \|X - Y\|$  for all  $X, Y \in \mathbb{R}^{m \times n}$ .*

Assumption 1(a) is standard. Assumption 1(b) is natural in the analysis of Muon-type algorithms (e.g., see (Chen et al., 2025; Riabinin et al., 2025; Shen et al., 2025)). It follows from Assumption 1(b) that

$$f(Y) \leq f(X) + \langle \nabla f(X), Y - X \rangle + \frac{L_*}{2} \|Y - X\|^2 \quad \forall X, Y \in \mathbb{R}^{m \times n}. \quad (4)$$

We next provide a definition for approximate stationary points of problem (1).

**Definition 1.** For any  $\epsilon \in (0, 1)$ , we say that a deterministic point  $X \in \mathbb{R}^{m \times n}$  is an  $\epsilon$ -nuclear norm stationary point (NSP) of problem (1) if it satisfies  $\|\nabla f(X)\|_* \leq \epsilon$ . We say that a point  $\hat{X} \in \mathbb{R}^{m \times n}$  is an  $\epsilon$ -stochastic nuclear norm stationary point (SNSP) of problem (1) if it satisfies  $\mathbb{E}[\|\nabla f(\hat{X})\|_*] \leq \epsilon$ , where the expectation is taken with respect to the randomness used by the stochastic algorithm (e.g., sampled gradients).

### 3 Matrix Optimization with Low-Rank Orthogonalization

In this section, we propose matrix optimization algorithms with low-rank orthogonalization for solving (1). In particular, we first propose low-rank orthogonalization in Section 3.1. Then, we propose low-rank MSGD in Section 3.2, and a low-rank variant of Muon in Section 3.3.

#### 3.1 Low-Rank Orthogonalization

Orthogonalization techniques have attracted increasing attention in recent optimizer designs, as it has shown strong empirical performance in foundation model training (e.g., see (Bernstein & Newhouse, 2024; Jordan et al.; Lau et al., 2025; Tuddenham et al., 2022)). In this subsection, we develop a low-rank orthogonalization method, leveraging low-rank matrix approximation techniques, that serves as a lightweight substitute for the existing orthogonalization subroutines used in matrix-variate optimizers.

---

##### Algorithm 1 A Low-Rank Orthogonalization Method

---

**Input:** matrix  $M \in \mathbb{R}^{m \times n}$ , rank trial  $r \in \mathbb{Z}_+ \cap [1, \varrho]$ .

**Output:** approximate matrix sign  $M_O \in \mathbb{R}^{m \times n}$ .

Draw a Gaussian random matrix  $G \in \mathbb{R}^{n \times r}$ .

Perform a QR decomposition on  $MG$  to obtain a column-orthogonal Q factor  $Q \in \mathbb{R}^{m \times r}$ .

Return  $M_O = Q \text{msgn}(Q^T M)$ . (On 1,  $\text{msgn}(Q^T M)$  is recommended to be estimated via Newton-Schulz iterations.)

---

Specifically, our low-rank orthogonalization method, presented in Algorithm 1, is based on Gaussian sketching (Halko et al., 2011). This method first draws a Gaussian random matrix  $G \in \mathbb{R}^{n \times r}$  with  $r \ll \varrho$ , and performs a QR decomposition on  $MG$  to obtain a column-orthogonal Q factor  $Q \in \mathbb{R}^{m \times r}$ . Then, it computes  $\text{msgn}(Q^T M) \in \mathbb{R}^{r \times n}$  and returns  $M_O = Q \text{msgn}(Q^T M)$  as a low-rank approximation for  $\text{msgn}(M)$ . As will be shown in Theorem 1,  $M_O$  represents the matrix sign of  $QQ^T M$ , which is a low-rank approximation of  $M$ . Its proof is deferred to Section 6.1.

**Theorem 1.** *Consider Algorithm 1 with inputs  $M \in \mathbb{R}^{m \times n}$  and  $r \in \mathbb{Z}_+ \cap [1, \varrho]$ , where  $\varrho := \min\{m, n\}$ . Let  $Q \in \mathbb{R}^{m \times r}$  be generated by Algorithm 1. Then, for any  $r_*$  satisfying  $2 \leq r_* \leq r - 2$ , it holds that*

$$\mathbb{E}[\|(I - QQ^T)M\|_F] \leq \left(1 + \frac{r_*}{r - r_* - 1}\right)^{1/2} \|M - [M]_{r_*}\|_F. \quad (5)$$

Moreover, we have  $\text{msgn}(QQ^T M) = Q \text{msgn}(Q^T M)$ .



**Remark 1.** The relation (5) is adapted from (Halko et al., 2011, Theorem 10.5), where additional guarantees—such as those involving different matrix norms and high-probability bounds—can also be found. In addition, low-rank matrix approximation based on column selection (e.g., see (Drineas et al., 2006a;b)) can be used to develop a low-rank orthogonalization method. However, since its approximation guarantee is more complicated than that of the Gaussian sketching-based approach, we defer the column-selection-based method to the supplementary materials.

Next, we illustrate two major advantages of our low-rank orthogonalization method, namely, *computational efficiency* and *noise robustness*, through synthetic experiments on randomly generated matrices.

**Computational Efficiency** We compare the computation time on GPUs for calculating inexact matrix sign of high-dimensional matrices using Newton-Schulz iterations, low-rank orthogonalization based on Gaussian sketching (Algorithm 1) and column selection (see supplementary materials), and truncated SVD.

For each  $n \in \{1000, 2000, 5000, 10000\}$ , we generate 50 random matrices  $M \in \mathbb{R}^{n \times n}$ , with each entry drawn from the standard Gaussian distribution. We then apply all competing orthogonalization methods to estimate  $\text{msgn}(M)$ . We implement Newton-Schulz iterations as provided in Muon (Jordan et al.) with 5 iterations. For our low-rank orthogonalization methods, we set  $r = 0.1n$  as the input rank parameter, use command `torch.linalg.qr` to compute the Q factor, and apply 5 iterations of Newton-Schulz scheme following Muon (Jordan et al.) to compute the matrix sign of  $Q^T M$ . In addition, we use command `torch.pca_lowrank` to efficiently perform truncated SVD, setting the rank to  $0.1n$  for truncation.




Figure 1: Left: Comparison of GPU computation time across Newton-Schulz iterations (NS), our low-rank orthogonalization with Gaussian sketching (GS) and column selection (CS), and truncated SVD (TSVD). Right: Speedup ratios of our low-rank orthogonalization methods compared to Newton-Schulz iterations.




Figure 2: Time distribution for low-rank orthogonalization with Gaussian sketching, including the QR decomposition, the Newton-Schulz iterations for computing the matrix sign, and other computations.

We present a comparison of the computation time on GPUs for all competing methods in Figure 1, and the time distribution of our low-rank orthogonalization method with Gaussian sketching in Figure 2. From Figure 1, we observe that our low-rank orthogonalization method significantly reduces computation time compared to Newton-Schulz iterations and truncated SVD. Although both truncated SVD and our low-rank orthogonalization exploit low-rank structure for computation, truncated SVD is not well-suited for GPU

environments and is therefore slower than the Newton-Schulz iterations. From Figure 2, we observe that in our low-rank orthogonalization method with Gaussian sketching, the QR decomposition accounts for approximately half of the total time, while the Newton-Schulz iterations and other operations (such as generating Gaussian random matrices and performing matrix multiplications) each make up roughly half of the remaining time.

**Noise Robustness** In addition to reducing computation time, our low-rank orthogonalization methods also produce more robust estimates of the matrix sign for low-rank matrices in the presence of noise. This robustness is particularly important in foundation model training, where gradients often exhibit low-rank structure. We now compare the performance of Newton-Schulz iterations and our low-rank methods in estimating the matrix sign of noisy, low-rank matrices.

For each  $n \in \{1000, 2000, 5000, 10000\}$ , we first randomly generate 10 nearly low-rank matrices  $M \in \mathbb{R}^{n \times n}$  as follows: the top  $0.1n$  singular values are set to 1, the remaining singular values are set to  $10^{-4}$ , and the singular vectors are randomly generated orthogonal vectors. For each  $M$ , we generate 50 noise matrices  $N \in \mathbb{R}^{n \times n}$ , with each entry drawn from a Gaussian distribution with mean zero and standard deviation  $\sigma \in \{10^{-3}, 10^{-2}, 10^{-1}, 5 \times 10^{-1}, 1\}$ , and construct noisy matrices  $M_N = M + N$ . We next apply Newton-Schulz iterations and Gaussian-sketching low-rank orthogonalization (GS), using an input rank parameter  $0.1n$ , to estimate  $\text{msgn}(M_N)$ . For each  $M$ , we vectorize each matrix-sign estimate, compute the empirical covariance matrix across the 50 noisy trials, and report its trace. This scalar equals the total coordinate-wise variance of the estimator, so a smaller value indicates a more stable matrix-sign estimate under noise. Both methods are implemented in the same way as in the computational-efficiency experiment above.

To complement this variance-based robustness metric, we also include ground-truth accuracy diagnostics using a clean matrix-sign reference. Specifically, for instances with  $n \in \{1000, 2000, 5000, 10000\}$ , we generate clean rank- $0.1n$  matrices  $M$  with nonzero singular values equal to 1000, compute the clean reference  $S_\star = \text{msgn}(M)$  from the generated low-rank factors, and then compare matrix-sign estimates  $\hat{S}(M_N)$  computed from noisy observations  $M_N = M + N$  against  $S_\star$ . We report the relative Frobenius error  $\|\hat{S}(M_N) - S_\star\|_F / \|S_\star\|_F$  and the alignment error  $1 - \langle \hat{S}(M_N), S_\star \rangle / (\|\hat{S}(M_N)\|_F \|S_\star\|_F)$ . For Gaussian-sketching (GS) low-rank orthogonalization, we test input ranks  $0.025n$ ,  $0.05n$ ,  $0.075n$ , and  $0.1n$ , which are smaller than or equal to the clean rank  $0.1n$ ; Newton-Schulz, which has no rank parameter, is applied to the full noisy matrix and plotted over the same rank range as a horizontal baseline. For each dimension, we generate one base matrix and one noisy realization with noise standard deviation  $\sigma = 0.5$ . These diagnostics distinguish estimator stability under noise from accuracy relative to the clean low-rank matrix-sign target.

Figure 3 shows that, when the target is the clean low-rank matrix sign, Gaussian-sketching low-rank orthogonalization yields substantially smaller relative Frobenius and alignment errors than applying Newton-Schulz iterations to the full noisy matrix. As the GS input rank increases from below the clean rank toward the clean rank  $0.1n$ , both errors decrease, showing that the approximation becomes more accurate as the rank captures more of the clean signal. Figure 4 presents the average total estimator variance, measured by the trace of the empirical covariance matrix, for each competing method. Compared to applying Newton-Schulz iterations to the full matrix, our low-rank orthogonalization method yields matrix-sign estimates that are substantially less sensitive to noise. Intuitively, singular vectors associated with very small singular values are the most unstable under perturbations, and truncating them removes the noisiest directions from the orthogonalization step. Taken together, Figures 3 and 4 show that low-rank orthogonalization improves robustness by filtering noise-dominated directions, while the choice of rank controls the tradeoff between noise filtering and retaining the clean low-rank signal. As a result, our low-rank method provides more stable estimates of the matrix sign, particularly for nearly low-rank matrices.

### 3.2 Low-Rank Matrix-Signed Gradient Descent

In this subsection, we propose low-rank MSGD methods, including a fixed-rank variant and a safeguarded variant with adaptive ranks. Both variants use low-rank orthogonalization as a subroutine.

At each iteration  $k \geq 0$ , the fixed-rank variant invokes Algorithm 1 to compute  $M_Q^k = \text{msgn}(M_Q^k)$ , where  $M_Q^k$  is a low-rank approximation of  $\nabla f(X^k)$ . Then, the next iterate  $X^{k+1}$  is obtained by performing a line search




Figure 3: Ground-truth accuracy diagnostics for recovering the clean low-rank matrix sign from noisy observations with noise standard deviation  $\sigma = 0.5$ . The four columns use  $n = 1000, 2000, 5000, 10000$ . The clean matrix has rank  $0.1n$ , and the horizontal axis gives the input rank used by Gaussian-sketching (GS) low-rank orthogonalization, ranging from below the clean rank to the clean rank. The rows report Frobenius error and alignment error with  $S_\star = \text{msgn}(M)$ ; the dashed horizontal lines show Newton-Schulz applied to the full noisy matrix as a rank-independent baseline.




Figure 4: Comparison of variance levels in matrix-sign estimation across Newton-Schulz iterations (NS) and Gaussian-sketching low-rank orthogonalization (GS), applied to nearly low-rank matrices under noise. GS uses input rank  $0.1n$ , matching the original RS rank setting. The horizontal axis reports the noise scale  $\sigma \in \{10^{-3}, 10^{-2}, 10^{-1}, 5 \times 10^{-1}, 1\}$  used to generate  $M_N = M + N$ .

---

#### Algorithm 2 Low-Rank MSGD

---

**Input:** starting point  $X^0 \in \mathbb{R}^{m \times n}$ , rank parameter  $r \in \mathbb{Z}_+ \cap [1, \varrho]$ , step sizes  $\{\eta_k\} \subset (0, \infty)$ .  
**for**  $k = 0, 1, 2, \dots$  **do**  
    Call Algorithm 1 (see Section 3.1) with  $(M, r) = (\nabla f(X^k), r)$  to obtain an approximate matrix sign  $M_O^k$ , such that  $M_O^k = \text{msgn}(M_Q^k)$ , where  $M_Q^k$  is a low-rank approximation for  $\nabla f(X^k)$ .  
    Update the next iterate:  $X^{k+1} = X^k - \eta_k M_O^k$ .  
**end for**

---

update from  $X^k$  along the matrix-signed direction  $-M_O^k$ . Details of this method are presented in Algorithm 2.

The following theorem provides a convergence guarantee for Algorithm 2, whose proof is deferred to Section 6.2.

**Theorem 2.** *Suppose that Assumption 1 holds. Let  $\{(X^k, M_O^k)\}$  be the sequence generated by Algorithm 2 with  $M_O^k = \text{msgn}(M_Q^k)$  for all  $k \geq 0$  and step sizes  $\{\eta_k\}$  given by  $\eta_k = (k+1)^{-1/2}$  for all  $k \geq 0$ . Then, it holds that for all  $K \geq 3$ ,*

$$\min_{0 \leq k \leq K-1} \{\|\nabla f(X^k)\|\} \leq \frac{f(X^0) - f_{\text{low}} + L_* \ln K}{K^{1/2}} + \frac{2}{K^{1/2}} \sum_{k=0}^{K-1} \frac{\|\nabla f(X^k) - M_Q^k\|_*}{(k+1)^{1/2}}, \quad (6)$$

where  $L_*$  is defined in Assumption 1.

**Remark 2.** From Theorem 2, we observe that when  $\{(X^k, M_Q^k)\}$  satisfies  $\sum_{k=0}^{K-1} \|\nabla f(X^k) - M_Q^k\|_*(k+1)^{-1/2} = \tilde{O}(1)$ , it holds that  $\min_{0 \leq k \leq K-1} \{\|\nabla f(X^k)\|_*\} = \tilde{O}(K^{-1/2})$ , which matches, up to a logarithmic factor, the well-established optimal convergence rate for nonconvex optimization (e.g., see (Carmon et al., 2020)). Moreover, if the gradients  $\{\nabla f(X^k)\}$  have low rank when  $k$  is sufficiently large (as is often the case during deep neural network training (Zhao et al., 2024)), we can tune the rank parameter  $r$  to be close to the effective rank of the gradients, causing the sequence  $\{\|\nabla f(X^k) - M_Q^k\|_*\}$  to remain close to zero. More concretely, if  $(\nabla f(X^k))_r$  denotes the best rank- $r$  approximation of  $\nabla f(X^k)$ , then  $\|\nabla f(X^k) - (\nabla f(X^k))_r\|_* = \sum_{i>r} \sigma_i(\nabla f(X^k))$ . Therefore, a sufficient per-iteration rank  $r_k$  is any rank for which this singular-value tail is below the target tolerance  $\delta_k$ ; an analogous Frobenius residual condition is  $\sum_{i>r_k} \sigma_i(\nabla f(X^k))^2 \leq \delta_k^2$ .

We next describe a safeguarded low-rank MSGD method with adaptively updated ranks. At each iteration  $k \geq 0$ , this method invokes Algorithm 1 with  $(M, r) = (\nabla f(X^k), r_k)$  to obtain  $M_O^k = \text{msgn}(M_Q^k)$ , where  $M_Q^k$  is a low-rank approximation of  $\nabla f(X^k)$  such that the approximation error satisfies (7). Then, the next iterate  $X^{k+1}$  is obtained by performing a line search update from  $X^k$  along the matrix-signed direction  $-M_O^k$ , with a suitable step size. Details of this method are provided in Algorithm 3.

It is noteworthy that the approximation error in (7) holds for some  $r_k \leq \varrho$  because when  $r_k = \varrho$ , the error  $\|M_Q^k - \nabla f(X^k)\|_*$  is zero. When the matrix is low-rank or has a small effective rank,  $r_k$  can be chosen to be much smaller than  $\varrho$ . In principle, one can gradually increase the trial ranks to find  $r_k$  such that (7) holds.

---

**Algorithm 3** Safeguarded Low-Rank MSGD

---

**Input:** starting point  $X^0 \in \mathbb{R}^{m \times n}$ , initial rank trial  $r_0 \in \mathbb{Z}_+ \cap [1, \varrho]$ , step sizes  $\{\eta_k\} \subset (0, \infty)$ , control errors  $\{\delta_k\} \subset (0, \infty)$

**for**  $k = 0, 1, 2, \dots$  **do**

    Call Algorithm 1 (see Section 3.1) with  $(M, r) = (\nabla f(X^k), r_k)$  to obtain an approximate matrix sign  $M_O^k$  such that  $M_O^k = \text{msgn}(M_Q^k)$  and

$$\|\nabla f(X^k) - M_Q^k\|_* \leq \delta_k. \quad (7)$$

    Update the next iterate:  $X^{k+1} = X^k - \eta_k M_O^k$ .

**end for**

---

The following theorem establishes an iteration complexity of Algorithm 3, whose proof is deferred to Section 6.2.

**Theorem 3.** *Suppose that Assumption 1 holds. Let  $f_{\text{low}}$  and  $L_*$  be given in Assumption 1, and define*

$$U_{\text{gd}} := f(X^0) - f_{\text{low}} + L_* + 4. \quad (8)$$

*Let  $\{X^k\}$  be generated by Algorithm 3 with inputs  $\{(\eta_k, \delta_k)\}$  given by  $\eta_k = \delta_k = (k+1)^{-1/2}$  for all  $k \geq 0$ . Then, for any  $\epsilon \in (0, 1)$ , it holds that  $\min_{0 \leq k \leq K-1} \{\|\nabla f(X^k)\|_*\} \leq \epsilon$  for all  $K$  satisfying*

$$K \geq \max \left\{ \left( \frac{4U_{\text{gd}}}{\epsilon} \ln \left( \frac{4U_{\text{gd}}}{\epsilon} \right) \right)^2, 3 \right\}.$$

**Remark 3.** From Theorem 3, we observe that Algorithm 3 achieves an iteration complexity of  $\tilde{O}(\epsilon^{-2})$  for finding an  $\epsilon$ -NSP of (1). This complexity bound matches, up to a polylogarithmic factor, the lower complexity bound as established in (Carmon et al., 2020).

### 3.3 Low-Rank Muon

In this subsection, we propose a low-rank variant of Muon (Jordan et al.), and analyze its iteration complexity under heavy-tailed noise.

This method follows a framework similar to Muon (Jordan et al.), but instead of directly computing the matrix sign of the momentum update, it computes only its low-rank approximation. Specifically, at each iteration of this method, it first performs a momentum update to generate  $M^k$  by aggregating stochastic gradients of  $f$  evaluated at  $X^0, \dots, X^k$ . Then, Algorithm 1 is invoked to obtain  $M_O^k = \text{msgn}(M_Q^k)$ , where  $M_Q^k$  is a low-rank approximation of  $\nabla f(X^k)$  such that the approximation error satisfies (10). The next iterate  $X^{k+1}$  is obtained by performing a line search update from  $X^k$  along the matrix-signed direction  $-M_O^k$  with a suitable step size. Details of this method are given in Algorithm 4.

Before analyzing the complexity of Algorithm 4 for computing an approximate solution to (1), we make the following heavy-tailed noise assumption regarding the stochastic gradient  $G(\cdot; \xi)$ .

**Assumption 2.** *The stochastic gradient estimator  $G : \mathbb{R}^{m \times n} \times \Xi \rightarrow \mathbb{R}^{m \times n}$  satisfies*

$$\mathbb{E}[G(X; \xi)] = \nabla f(X), \quad \mathbb{E}[\|G(X; \xi) - \nabla f(X)\|_F^\alpha] \leq \sigma^\alpha$$

for some  $\sigma > 0$  and  $\alpha \in (1, 2]$ .

---

#### Algorithm 4 Low-Rank Muon

---

**Input:** starting point  $X^0 \in \mathbb{R}^{m \times n}$ , initial rank trial  $r_0 \in \mathbb{Z}_+ \cap [1, \varrho]$ , step sizes  $\{\eta_k\} \subset (0, \infty)$ , weighting parameters  $\{\theta_k\} \subset (0, 1]$ , control errors  $\{\delta_k\} \subset (0, \infty)$ .

**Initialize:**  $M^{-1} = \mathbf{0}_{m \times n}$  and  $\theta_{-1} = 1$ .

**for**  $k = 0, 1, 2, \dots$  **do**

Compute the full-rank search direction:

$$M^k = (1 - \theta_{k-1})M^{k-1} + \theta_{k-1}G(X^k; \xi^k). \quad (9)$$

Call Algorithm 1 (see Section 3.1) with  $(M, r) = (M^k, r_k)$  to obtain an approximate matrix sign  $M_O^k$  such that  $M_O^k = \text{msgn}(M_Q^k)$  and

$$\|M^k - M_Q^k\|_* \leq \delta_k. \quad (10)$$

Update the next iterate:

$$X^{k+1} = X^k - \eta_k M_O^k. \quad (11)$$

**end for**

---

The following theorem establishes the iteration complexity of Algorithm 4, whose proof is deferred to Section 6.3.

**Theorem 4.** *Suppose that Assumptions 1 and 2 hold. Let  $f_{\text{low}}$  and  $L_*$  be given in Assumption 1, and  $\sigma$  and  $\alpha$  be given in Assumption 2. Define*

$$U_{\text{mn}} := f(X^0) - f_{\text{low}} + \sigma^\alpha + 2L_* + 4 + 2(\alpha - 1)(2\varrho^{1/2}/\alpha)^{\alpha/(\alpha-1)} + 6L_*^\alpha + 4\sigma^\alpha. \quad (12)$$

*Let  $\{X^k\}$  be generated by Algorithm 4 with input parameters  $\{(\eta_k, \theta_k, \delta_k)\}$  given by  $\eta_k = (k+1)^{-(2\alpha-1)/(3\alpha-2)}$ ,  $\theta_k = (k+1)^{-\alpha/(3\alpha-2)}$ , and  $\delta_k = (k+1)^{-(\alpha-1)/(3\alpha-2)}$  for all  $k \geq 0$ . Then, for any  $\epsilon \in (0, 1)$ , it holds that  $\mathbb{E}[\|\nabla f(X^{\iota_K})\|_*] \leq \epsilon$  for all  $K$  satisfying*

$$K \geq \max \left\{ \left( \frac{2(3\alpha-2)U_{\text{mn}}}{(\alpha-1)\epsilon} \ln \left( \frac{2(3\alpha-2)U_{\text{mn}}}{(\alpha-1)\epsilon} \right) \right)^{\frac{3\alpha-2}{\alpha-1}}, 3 \right\},$$

where  $\iota_K$  is uniformly drawn from  $\{0, \dots, K-1\}$ .

**Remark 4.** From Theorem 4, one can observe that Algorithm 4 achieves an iteration complexity of  $\tilde{\mathcal{O}}(\epsilon^{-(3\alpha-2)/(\alpha-1)})$  for finding an  $\epsilon$ -SNSP of problem (1), which matches the optimal dependence on  $\epsilon$  in the complexity results for vector-variate stochastic first-order methods under heavy-tailed noise (see, e.g., (He et al., 2025; Zhang et al., 2020)). To the best of our knowledge, our result is the first to show that inexact Muon-type algorithms can achieve an iteration complexity with optimal dependence on  $\epsilon$  when applied to matrix-variate optimization under heavy-tailed noise. For the low-rank Muon setting, the same rank requirement should be read in terms of the singular-value tail of the momentum matrix  $M^k$  being orthogonalized: if  $(M^k)_r$  is its best rank- $r$  approximation, then choosing  $r_k$  so that  $\sum_{i>r_k} \sigma_i(M^k)$ , or the corresponding Frobenius tail, is below the prescribed inexactness level is sufficient. Thus the theorem assumes a low-effective-rank condition along the trajectory rather than guaranteeing that a uniformly small rank is always enough. We emphasize again that this statement applies to the residual-controlled inexact algorithm in Algorithm 4; our fixed-rank practical implementation is a computationally motivated specialization.

## 4 Numerical Experiments

In this section, we present numerical experiments to evaluate the performance of low-rank Muon and compare our method with vanilla Muon, AdamW, and SGD. The experiments are conducted on Matrix Regression (Section 4.1), GPT-2 pretraining (Section 4.2) and LLaMA pretraining (Section 4.3) from random initialization, without using pretrained checkpoints. All experiments are executed on a server with 2 NVIDIA A100 GPUs (80 GB each). To test the effectiveness of Low-Rank Muon in Algorithm 4, we first use fixed ranks 100 and 200, and also report AdaRank runs using the lightweight adaptive-rank implementation described below.

In practice, AdaRank chooses the rank once per optimizer step for each two-dimensional momentum matrix. Let  $G_k$  be the current gradient matrix and

$$M_k = \mu M_{k-1} + (1 - \mu)G_k$$

be the momentum matrix to be orthogonalized. The implementation does not compute a full SVD and does not scan many trial ranks. Instead, it estimates a stable-rank proxy using only two power iterations. For a matrix  $M_k \in \mathbb{R}^{m \times n}$ , let  $\bar{r} = \min\{r_{\max}, m, n\}$ . The power iterations give a unit vector  $v_k$  and the spectral-norm estimate  $\hat{\sigma}_{1,k}^2 = \|M_k v_k\|_2^2$ . With the default settings  $r_0 = 64$ , rank step  $\Delta r = 32$ , maximum rank  $r_{\max} = 256$ ,  $\delta_0 = 0.1$ , and  $\gamma = 0.25$ , we set  $\delta_k = \delta_0(k+1)^{-\gamma}$  and compute

$$\hat{s}_k = \text{clip} \left( \frac{\|M_k\|_F^2}{\hat{\sigma}_{1,k}^2 + 10^{-12}}, 1, \bar{r} \right), \quad c_k = 1.25 + \min\{0.50, 0.10 \max(0, -\log_{10} \delta_k)\},$$

and choose

$$r_k = \min \left\{ \bar{r}, \max \left\{ r_0, \Delta r \left\lceil \frac{c_k \hat{s}_k}{\Delta r} \right\rceil \right\} \right\}.$$

Thus the selected rank is a rounded and clipped multiple of the estimated stable rank. The reason for using stable rank is that  $\|M_k\|_F^2 / \|M_k\|_2^2$  measures the effective number of dominant singular directions: it is small when the spectrum is concentrated and increases when more directions carry comparable energy. It can also be estimated cheaply from a Frobenius norm and a spectral-norm estimate, so it gives a practical rank signal without paying for an SVD. The safety factor and upward rounding reduce underestimation, while the cap  $r_{\max}$  controls optimizer overhead. After selecting  $r_k$ , the implementation draws one randomized column projection: it samples  $r_k$  columns, forms  $Y_k = M_k(:, \Omega_k)$ , computes a reduced QR factorization  $Y_k = Q_k R_k$ , sets  $B_k = Q_k^T M_k$ , and applies Newton–Schulz orthogonalization to the smaller matrix  $B_k$ . The resulting low-rank matrix-sign estimate is  $\hat{S}_k = Q_k \text{NS}(B_k)$ , with a transpose-and-transpose-back convention for tall matrices. The parameter is then updated by the usual Muon-style step using  $\hat{S}_k$ . This implementation increases the rank when the estimated stable rank grows, while each matrix uses only one projection per step.

As additional empirical context from the completed training logs, we report completed-run efficiency diagnostics in Tables 4 and 5 of Appendix 5. These diagnostics summarize the completed full-budget runs together with the read-off curve and timing/profile readouts used for the efficiency comparison.

## 4.1 Matrix Regression

In this subsection, we consider the synthetic matrix regression problem

$$\min_{X \in \mathbb{R}^{n \times n}} \left\{ f(X) = \frac{1}{2} \|AXB - C\|_F^2 \right\}. \quad (13)$$

Here  $A \in \mathbb{R}^{p \times n}$ ,  $B \in \mathbb{R}^{n \times p}$ , and  $C \in \mathbb{R}^{p \times p}$  are given data matrices. We simulate low-rank structure for the gradients  $\nabla f(X)$  by choosing  $n \gg p = 100$ . For each  $n \in \{10^3, 5 \times 10^3, 10^4\}$ , we set  $A = U_A \Sigma_A V_A^T$  and  $B = U_B \Sigma_B V_B^T$ , where  $\Sigma_A = \Sigma_B = \text{Diag}(2^{-1}, \dots, 2^{-p})$ , and  $U_A, V_A, U_B$ , and  $V_B$  are randomly generated column-orthogonal matrices. We also generate a ground truth solution  $X^*$  with all entries independently drawn from the standard Gaussian distribution, and set  $C = AX^*B + 10^{-3}E$ , where all entries of  $E$  are independently drawn from the standard Gaussian distribution.




Figure 5: Objective values (top row) and gradient Frobenius norms (bottom row) for all methods during the first  $2 \times 10^4$  iterations. Results for  $n = 10^3$ ,  $5 \times 10^3$ , and  $10^4$  are shown in the left, middle, and right columns, respectively. The computation times on GPUs are displayed in the legend for each plot.

We apply low-rank Muon, Muon, AdamW, and SGD to solve (13). We implement two fixed-rank low-rank Muon variants with rank parameters 50 and 100, abbreviated as LR-Muon50 and LR-Muon100, respectively. We also include AdaRank, which uses the adaptive-rank rule described at the beginning of Section 4; in this deterministic full-batch problem, the gradient entering the momentum update is  $G_k = \nabla f(X^k) = A^T(AX^k B - C)B^T$ . All methods are initialized from the zero matrix. We compare their performance based on the objective values and the Frobenius norm of the gradient for problem (13), evaluated over the first  $2 \times 10^4$  iterations. The algorithmic parameters are selected to suit each method well in terms of computational performance.

For each  $n$ , we plot the objective values and the gradient Frobenius norms in Figure 5 to illustrate the convergence behavior of all methods. Figure 5 shows that our fixed-rank and adaptive-rank low-rank Muon variants achieve convergence behavior comparable to vanilla Muon in iteration count, while clearly outperforming SGD and AdamW on this synthetic matrix problem. In particular, AdaRank tracks the fixed-rank low-rank Muon curves closely while adaptively adjusting  $r_k$  along the trajectory, which supports the rank-adaptation mechanism in a controlled setting. In addition, SGD fails to converge effectively for problem (13). The low-rank variants also exhibit lower GPU computation time than vanilla Muon in these synthetic experiments, which is consistent with the efficiency motivation of the proposed low-rank orthogonalization. Since this experiment uses deterministic full-batch gradients and a controlled low-rank structure, we view it as a complementary sanity check.

## 4.2 GPT-2 Pretraining

In this subsection, we consider GPT-2 (Radford et al., 2019) pretraining from random initialization. We experiment with GPT-2 models of sizes 60M, 135M, 350M, and 1B on the same three datasets tested in the Muon GitHub repository (Jordan et al.): FineWeb10B, FineWeb100B, and FineWebEdu10B. No pretrained checkpoints are used in these runs.

Table 1: Comparison of validation perplexity and computational time for all competing methods on GPT-2 pretraining. The  $\pm$  values report empirical variability over three runs with different random seeds.

| Dataset       | Method     | Validation Perplexity   |                         |                         |                         | Computational Time |        |        |        |
|---------------|------------|-------------------------|-------------------------|-------------------------|-------------------------|--------------------|--------|--------|--------|
|               |            | 60M                     | 135M                    | 350M                    | 1B                      | 60M                | 135M   | 350M   | 1B     |
| FineWeb10B    | SGDM       | 48.85 $\pm$ 0.44        | 77.50 $\pm$ 0.57        | 31.36 $\pm$ 0.33        | 27.86 $\pm$ 0.29        | 2.51e3             | 7.08e3 | 1.20e4 | 2.70e5 |
|               | AdamW      | 43.56 $\pm$ 0.35        | 43.56 $\pm$ 0.32        | 25.72 $\pm$ 0.27        | 21.58 $\pm$ 0.23        | 2.52e3             | 7.08e3 | 1.21e4 | 2.69e5 |
|               | GaLore     | 57.61 $\pm$ 0.49        | 39.59 $\pm$ 0.37        | 30.87 $\pm$ 0.31        | 28.58 $\pm$ 0.28        | 2.73e3             | 7.12e3 | 1.24e4 | 2.76e5 |
|               | SOAP       | 41.32 $\pm$ 0.38        | 55.37 $\pm$ 0.46        | 38.45 $\pm$ 0.34        | 34.58 $\pm$ 0.30        | 2.56e3             | 7.09e3 | 1.23e4 | 2.73e5 |
|               | Muon       | <b>32.89</b> $\pm$ 0.27 | 32.99 $\pm$ 0.26        | 28.48 $\pm$ 0.24        | 20.99 $\pm$ 0.21        | 2.76e3             | 7.57e3 | 1.32e4 | 3.04e5 |
|               | LR-Muon100 | 35.21 $\pm$ 0.30        | 28.83 $\pm$ 0.25        | 27.32 $\pm$ 0.23        | 21.70 $\pm$ 0.22        | 2.69e3             | 7.49e3 | 1.25e4 | 2.90e5 |
|               | LR-Muon200 | 33.98 $\pm$ 0.28        | <b>27.34</b> $\pm$ 0.24 | <b>25.64</b> $\pm$ 0.22 | <b>20.69</b> $\pm$ 0.20 | 2.79e3             | 7.54e3 | 1.26e4 | 2.91e5 |
|               | AdaRank    | 35.61 $\pm$ 0.31        | 28.31 $\pm$ 0.27        | 26.75 $\pm$ 0.24        | 21.45 $\pm$ 0.22        | 2.72e3             | 7.55e3 | 1.25e4 | 2.91e5 |
| FineWeb100B   | SGDM       | 51.48 $\pm$ 0.47        | 51.48 $\pm$ 0.43        | 31.54 $\pm$ 0.32        | 29.75 $\pm$ 0.31        | 2.69e3             | 7.25e3 | 1.04e4 | 2.52e5 |
|               | AdamW      | 43.52 $\pm$ 0.36        | 43.51 $\pm$ 0.33        | 27.47 $\pm$ 0.28        | 23.53 $\pm$ 0.24        | 2.70e3             | 7.21e3 | 1.05e4 | 2.56e5 |
|               | GaLore     | 56.42 $\pm$ 0.48        | 58.09 $\pm$ 0.50        | 31.19 $\pm$ 0.32        | 26.53 $\pm$ 0.27        | 2.81e3             | 7.31e3 | 1.15e4 | 2.67e5 |
|               | SOAP       | 41.32 $\pm$ 0.37        | 44.32 $\pm$ 0.39        | 30.98 $\pm$ 0.29        | 25.56 $\pm$ 0.25        | 2.71e3             | 7.23e3 | 1.07e4 | 2.58e5 |
|               | Muon       | 34.04 $\pm$ 0.29        | 33.02 $\pm$ 0.27        | 32.76 $\pm$ 0.26        | 20.75 $\pm$ 0.20        | 2.79e3             | 7.70e3 | 1.19e4 | 3.09e5 |
|               | LR-Muon100 | 35.78 $\pm$ 0.31        | 30.19 $\pm$ 0.26        | 27.59 $\pm$ 0.24        | 21.76 $\pm$ 0.21        | 2.77e3             | 7.62e3 | 1.09e4 | 2.87e5 |
|               | LR-Muon200 | <b>34.02</b> $\pm$ 0.29 | <b>28.31</b> $\pm$ 0.25 | <b>25.86</b> $\pm$ 0.23 | <b>20.45</b> $\pm$ 0.19 | 2.80e3             | 7.69e3 | 1.10e4 | 2.91e5 |
|               | AdaRank    | 35.93 $\pm$ 0.32        | 29.58 $\pm$ 0.28        | 26.76 $\pm$ 0.30        | 21.02 $\pm$ 0.23        | 2.79e3             | 7.68e3 | 1.08e4 | 2.89e5 |
| FineWebEdu10B | SGDM       | 50.92 $\pm$ 0.45        | 32.79 $\pm$ 0.33        | 31.54 $\pm$ 0.32        | 25.49 $\pm$ 0.26        | 2.56e3             | 7.09e3 | 1.20e4 | 2.53e5 |
|               | AdamW      | 43.56 $\pm$ 0.34        | 26.45 $\pm$ 0.27        | 21.01 $\pm$ 0.22        | 22.45 $\pm$ 0.24        | 2.59e3             | 7.09e3 | 1.20e4 | 2.56e5 |
|               | GaLore     | 69.49 $\pm$ 0.53        | 39.59 $\pm$ 0.36        | 57.42 $\pm$ 0.47        | 37.45 $\pm$ 0.35        | 2.64e3             | 7.12e3 | 1.24e4 | 2.60e5 |
|               | SOAP       | 55.37 $\pm$ 0.45        | 41.77 $\pm$ 0.38        | 47.90 $\pm$ 0.41        | 39.20 $\pm$ 0.36        | 2.60e3             | 7.11e3 | 1.22e4 | 2.58e5 |
|               | Muon       | 32.88 $\pm$ 0.28        | 23.89 $\pm$ 0.23        | 22.23 $\pm$ 0.22        | 19.54 $\pm$ 0.19        | 2.80e3             | 7.63e3 | 1.33e4 | 3.05e5 |
|               | LR-Muon100 | 35.60 $\pm$ 0.30        | 23.50 $\pm$ 0.22        | 22.39 $\pm$ 0.21        | 19.97 $\pm$ 0.20        | 2.74e3             | 7.53e3 | 1.26e4 | 2.85e5 |
|               | LR-Muon200 | 33.93 $\pm$ 0.27        | <b>22.37</b> $\pm$ 0.21 | <b>20.96</b> $\pm$ 0.20 | <b>18.85</b> $\pm$ 0.18 | 2.83e3             | 7.59e3 | 1.29e4 | 2.87e5 |
|               | AdaRank    | <b>32.51</b> $\pm$ 0.26 | 23.20 $\pm$ 0.23        | 21.20 $\pm$ 0.22        | 18.94 $\pm$ 0.21        | 2.73e3             | 7.58e3 | 1.26e4 | 2.84e5 |

We apply low-rank Muon, Muon, SOAP, GaLore, AdamW, and SGD with momentum (SGDM) for GPT-2 pretraining. We implement two practical low-rank Muon variants with fixed rank parameters 100 and 200, abbreviated as LR-Muon100 and LR-Muon200, respectively. We also report completed AdaRank runs using the adaptive-rank procedure defined at the beginning of Section 4; in these runs, the rank is updated once per optimizer step for each hidden matrix parameter, and the selected-rank statistics are logged during training. Following the experiments in (Jordan et al., 2024; Lau et al., 2025), we use AdamW to train the embedding and head layers for both Muon and low-rank Muon. All methods are trained from random initialization and terminated after one training epoch, consisting of 5000 iterations. We compare all methods using validation perplexity, a standard metric in language-model training (e.g., (Radford et al., 2019; Touvron et al., 2023; Vaswani et al., 2017)), which is defined as the exponential of the validation loss and serves to amplify performance differences. The algorithmic parameters are tuned for computational performance. For the adaptive-rank rows, we report the runs that completed the full training budget. In addition to final validation perplexity and wall-clock time, Appendix 5 reports how we compute or log time-to-target perplexity, tokens/sec, optimizer overhead, and peak memory. These auxiliary diagnostics should be interpreted as completed-run efficiency evidence rather than exhaustive rank-ablation estimates. More details about the experimental setups, including algorithm and GPT-2 model parameters, are provided in Appendix 5.

To further connect the practical implementation with the approximation conditions in Algorithms 3 and 4, Figure 6 visualizes the evolution of the leading singular values of representative Q, K, and V gradient matrices during GPT-2 pretraining. The spectra remain clearly decayed across iterations, which indicates that a moderate rank can capture a large portion of the update energy for substantial portions of training.



This qualitative behavior supports the low-rank motivation behind both the fixed-rank experiments and the residual-controlled rank-selection rule studied in the theory.




Figure 6: Evolution of the leading singular values of the Q, K, and V gradient matrices over training iterations for a representative GPT-2 350M pretraining run on FineWeb100B using Muon. The plot corresponds to a representative attention layer and illustrates persistent spectral decay throughout training.

Figure 7 gives the corresponding tail-residual diagnostic for GPT-2 large pretraining over the first 5000 optimizer iterations. We compute singular-value nuclear-tail residuals for ranks  $r = 100, 200, 400$ . The estimated tails decrease over training and are consistently smaller for larger ranks, providing additional empirical support that the sufficient ranks required by the approximation condition can remain moderate in the GPT-2 pretraining setting.




Figure 7: ESingular-value nuclear-tail residuals of representative Q, K, and V optimizer matrices in a GPT-2 large pretraining run over the first 5000 optimizer iterations. The colored curves show smoothed trends for ranks  $r = 100, 200, 400$ . The gray dashed curve is a normalized  $(k + 1)^{-1/2}$  reference, matching the approximation tolerance rate used in Theorem 3.

We present a comparison of validation perplexity and computational time in Table 1. For GPT-2 with a model size of 60M, our low-rank Muon achieves slightly worse validation perplexity than full-rank Muon, but still performs better than AdamW, SGDM, SOAP, and GaLore. For GPT-2 models with larger sizes, the fixed-rank low-rank Muon variants often improve upon vanilla Muon and the other baselines in terms of validation perplexity. We also observe that the computational time of our low-rank Muon is usually lower than that of vanilla Muon.

We plot the convergence behavior of validation perplexity versus training iterations in Figure 8. The figure shows similar qualitative trends across the three datasets. For the model size of 60M, all Muon-type methods consistently outperform AdamW and SGDM, while our low-rank Muon performs slightly worse than vanilla Muon. For the 135M and 350M models, the low-rank variants are often stronger than vanilla Muon. Taken together with Table 1, these results indicate that the fixed-rank low-rank approximation can be beneficial in GPT-2 pretraining, especially at larger model sizes.




Figure 8: Comparison of validation perplexity versus computational time for all competing methods in pretraining GPT-2 models.

### 4.3 LLaMA Pretraining

In this subsection, we consider LLaMA (Touvron et al., 2023) pretraining from random initialization. We experiment with LLaMA models of sizes 60M, 135M, 350M, and 1B on the same three datasets tested in the Muon GitHub repository (Jordan et al.): FineWeb10B, FineWeb100B, and FineWebEdu10B. No pretrained checkpoints are used in these runs.

We apply low-rank Muon, Muon, AdamW, GaLore, SOAP and SGD with momentum (SGDM) for LLaMA pretraining. Similar to Section 4.2, we implement two practical low-rank Muon variants with fixed rank parameters 100 and 200, and we use AdamW to train the embedding and head layers for both Muon and low-rank Muon. We also report completed AdaRank runs using the same adaptive-rank procedure defined at the beginning of Section 4, so the rank is updated once per optimizer step for each hidden matrix parameter before the low-rank orthogonalization is applied. All methods are trained from random initialization and terminated after one training epoch, consisting of 5000 iterations. We compare all methods using validation perplexity. The algorithmic parameters are tuned for computational performance. More details about the experimental setups, including algorithm and LLaMA model parameters, are provided in Appendix 5.

Figure 10 further visualizes the evolution of the leading singular values of representative Q, K, and V gradient matrices during LLaMA pretraining. Similar to the GPT-2 diagnostics, the spectra exhibit clear decay across iterations, supporting the empirical low-effective-rank structure used by the low-rank orthogonalization step.

Table 2: Comparison of validation perplexity and computational time for all competing methods on LLaMA pretraining. The  $\pm$  values report empirical variability over three runs with different random seeds.

| Dataset       | Method     | Validation Perplexity   |                         |                         |                         | Computational Time |        |        |        |
|---------------|------------|-------------------------|-------------------------|-------------------------|-------------------------|--------------------|--------|--------|--------|
|               |            | 60M                     | 135M                    | 350M                    | 1B                      | 60M                | 135M   | 350M   | 1B     |
| FineWeb10B    | SGDM       | 94.90 $\pm$ 0.64        | 85.99 $\pm$ 0.59        | 86.49 $\pm$ 0.58        | 44.71 $\pm$ 0.38        | 2.49e3             | 5.27e3 | 1.20e4 | 2.77e5 |
|               | AdamW      | 50.51 $\pm$ 0.42        | 40.17 $\pm$ 0.35        | 35.76 $\pm$ 0.31        | 24.14 $\pm$ 0.25        | 2.46e3             | 5.27e3 | 1.21e4 | 2.78e5 |
|               | GaLore     | 53.69 $\pm$ 0.68        | 47.70 $\pm$ 0.51        | 38.12 $\pm$ 0.52        | 26.54 $\pm$ 0.30        | 2.37e3             | 6.58e3 | 1.18e4 | 2.75e5 |
|               | SOAP       | 49.23 $\pm$ 0.45        | 50.23 $\pm$ 0.35        | 41.73 $\pm$ 1.08        | 26.34 $\pm$ 0.28        | 2.40e3             | 5.67e3 | 1.11e4 | 2.67e5 |
|               | Muon       | 36.68 $\pm$ 0.31        | 37.96 $\pm$ 0.32        | 36.43 $\pm$ 0.30        | 20.99 $\pm$ 0.21        | 2.59e3             | 5.69e3 | 1.35e4 | 3.18e5 |
|               | LR-Muon100 | 37.33 $\pm$ 0.32        | <b>34.08</b> $\pm$ 0.28 | 33.54 $\pm$ 0.27        | 21.70 $\pm$ 0.22        | 2.69e3             | 5.66e3 | 1.28e4 | 2.79e5 |
|               | LR-Muon200 | <b>36.05</b> $\pm$ 0.30 | 34.37 $\pm$ 0.29        | <b>32.37</b> $\pm$ 0.27 | <b>20.69</b> $\pm$ 0.20 | 2.70e3             | 5.70e3 | 1.30e4 | 2.91e5 |
|               | AdaRank    | 38.15 $\pm$ 0.37        | 34.96 $\pm$ 0.34        | 33.80 $\pm$ 0.30        | 22.98 $\pm$ 0.22        | 2.69e3             | 5.69e3 | 1.29e4 | 2.90e5 |
| FineWeb100B   | SGDM       | 81.59 $\pm$ 0.55        | 77.50 $\pm$ 0.52        | 87.49 $\pm$ 0.60        | 37.56 $\pm$ 0.34        | 2.60e3             | 5.44e3 | 1.05e4 | 2.65e5 |
|               | AdamW      | 41.49 $\pm$ 0.36        | 40.54 $\pm$ 0.34        | 35.47 $\pm$ 0.30        | 23.89 $\pm$ 0.24        | 2.61e3             | 5.47e3 | 1.06e4 | 2.69e5 |
|               | GaLore     | 58.97 $\pm$ 0.46        | 45.14 $\pm$ 0.38        | 37.05 $\pm$ 0.33        | 35.56 $\pm$ 0.31        | 2.65e3             | 5.50e3 | 1.09e4 | 2.69e5 |
|               | SOAP       | 55.08 $\pm$ 0.44        | 52.28 $\pm$ 0.42        | 47.52 $\pm$ 0.39        | 43.89 $\pm$ 0.36        | 2.64e3             | 5.50e3 | 1.07e4 | 2.70e5 |
|               | Muon       | 38.49 $\pm$ 0.33        | 37.66 $\pm$ 0.31        | 36.66 $\pm$ 0.30        | 21.76 $\pm$ 0.22        | 2.58e3             | 5.63e3 | 1.07e4 | 3.16e5 |
|               | LR-Muon100 | 37.68 $\pm$ 0.32        | 34.58 $\pm$ 0.29        | 33.71 $\pm$ 0.28        | 22.86 $\pm$ 0.23        | 2.63e3             | 5.55e3 | 1.10e4 | 2.79e5 |
|               | LR-Muon200 | 36.26 $\pm$ 0.31        | <b>34.20</b> $\pm$ 0.28 | <b>32.73</b> $\pm$ 0.27 | <b>20.98</b> $\pm$ 0.21 | 2.75e3             | 5.69e3 | 1.11e4 | 2.88e5 |
|               | AdaRank    | <b>34.90</b> $\pm$ 0.36 | 34.56 $\pm$ 0.33        | 32.89 $\pm$ 0.45        | 21.89 $\pm$ 0.30        | 2.77e3             | 5.78e3 | 1.11e4 | 2.82e5 |
| FineWebEdu10B | SGDM       | 82.63 $\pm$ 0.56        | 73.48 $\pm$ 0.50        | 73.48 $\pm$ 0.49        | 47.56 $\pm$ 0.39        | 2.57e3             | 5.27e3 | 1.20e4 | 2.76e5 |
|               | AdamW      | 41.82 $\pm$ 0.35        | 32.96 $\pm$ 0.29        | 29.12 $\pm$ 0.26        | 22.54 $\pm$ 0.23        | 2.60e3             | 5.29e3 | 1.20e4 | 2.78e5 |
|               | GaLore     | 64.04 $\pm$ 0.49        | 68.13 $\pm$ 0.52        | 45.13 $\pm$ 0.37        | 43.56 $\pm$ 0.35        | 2.57e3             | 5.27e3 | 1.20e4 | 2.76e5 |
|               | SOAP       | 40.32 $\pm$ 0.34        | 28.10 $\pm$ 0.27        | 27.79 $\pm$ 0.26        | 24.54 $\pm$ 0.24        | 2.60e3             | 5.29e3 | 1.20e4 | 2.78e5 |
|               | Muon       | <b>28.57</b> $\pm$ 0.25 | 30.72 $\pm$ 0.27        | 29.47 $\pm$ 0.26        | 22.67 $\pm$ 0.22        | 2.68e3             | 5.84e3 | 1.33e4 | 3.25e5 |
|               | LR-Muon100 | 30.52 $\pm$ 0.27        | 27.72 $\pm$ 0.25        | 27.24 $\pm$ 0.24        | 21.36 $\pm$ 0.21        | 2.72e3             | 5.66e3 | 1.24e4 | 2.88e5 |
|               | LR-Muon200 | 29.29 $\pm$ 0.26        | <b>27.65</b> $\pm$ 0.24 | <b>26.87</b> $\pm$ 0.23 | <b>19.54</b> $\pm$ 0.19 | 2.81e3             | 5.70e3 | 1.26e4 | 2.98e5 |
|               | AdaRank    | 32.30 $\pm$ 0.30        | 29.51 $\pm$ 0.29        | 28.35 $\pm$ 0.27        | 23.45 $\pm$ 0.20        | 2.83e3             | 5.68e3 | 1.27e4 | 2.89e5 |

To make the rank requirement more explicit, Figure 11 plots singular-value nuclear-tail residuals for representative Q, K, and V optimizer matrices during the first 5000 optimizer iterations of a LLaMA large pretraining run. We compute the nuclear-tail residuals for several ranks to evaluate how the singular-value tail changes over training. The estimated tails decrease over training and are smaller for larger ranks, which is consistent with the sufficient-rank interpretation of the approximation assumptions in Sections 3.2 and 3.3.

We present a comparison of validation perplexity and computational time in Table 2. Across most tested model sizes and datasets, the low-rank Muon variants achieve better validation perplexity than Muon, GaLore, and SOAP, while also clearly outperforming AdamW and SGDM in perplexity. We also observe that the computational time of low-rank Muon is often lower than that of vanilla Muon. These observations suggest that low-rank orthogonalization is a promising practical alternative to full-rank Newton-Schulz orthogonalization in these LLaMA pretraining settings.

We plot the convergence behavior of validation perplexity versus training iterations in Figure 9. For the model size of 60M, all Muon-type methods substantially outperform AdamW and SGDM, while low-rank Muon is similar to vanilla Muon. For the 135M and 350M models, the low-rank variants are generally stronger than vanilla Muon and the other baselines. Overall, these curves support the claim that low-rank Muon is competitive and often advantageous in the tested LLaMA pretraining settings, while the gains remain less uniform at the smallest scale.




Figure 9: Comparison of validation perplexity and computational time for all competing methods in pretraining LLaMA models.




Figure 10: Evolution of the leading singular values of the Q, K, and V gradient matrices over training iterations for a representative LLaMA pretraining run. The plot corresponds to a representative attention layer and illustrates persistent spectral decay throughout training.

## References

- Kwangjun Ahn and Byron Xu. Dion: A communication-efficient optimizer for large models. *arXiv preprint arXiv:2504.05295*, 2025.
- Kwangjun Ahn, Byron Xu, Natalie Abreu, and John Langford. Dion: Distributed orthonormalized updates. *arXiv preprint arXiv:2504.05295*, 2025.




Figure 11: Singular-value nuclear-tail residuals of representative Q, K, and V optimizer matrices in a LLaMA large pretraining run over the first 5000 optimizer iterations. The colored curves show smoothed trends for ranks  $r = 100, 200, 400$ . The gray dashed curve is a normalized  $(k + 1)^{-1/2}$  reference, matching the approximation tolerance rate used in Theorem 3.

Essential AI, Ishaan Shah, Anthony M. Polloreno, Karl Stratos, Philip Monk, Adarsh Chaluvvaraju, Andrew Hojel, Andrew Ma, Anil Thomas, Ashish Tanwer, Darsh J. Shah, Khoi Nguyen, Kurt Smith, Michael Callahan, Michael Pust, Mohit Parmar, Peter Rushton, Platon Mazarakis, Ritvik Kapila, Saurabh Srivastava, Somanshu Singla, Tim Romanski, Yash Vanjani, and Ashish Vaswani. Practical efficiency of Muon for pretraining. *arXiv preprint arXiv:2505.02222*, 2025.

Kang An, Yuxing Liu, Rui Pan, Shiqian Ma, Donald Goldfarb, and Tong Zhang. ASGO: Adaptive structured gradient optimization. *arXiv preprint arXiv:2503.20762*, March 2025. doi: 10.48550/arXiv.2503.20762. Preprint.

Rohan Anil, Vineet Gupta, Tomer Koren, Kevin Regan, and Yoram Singer. Scalable second order optimization for deep learning. *arXiv preprint arXiv:2002.09018*, 2020.

Jeremy Bernstein and Laker Newhouse. Old optimizer, new norm: An anthology. In *Workshop on Optimization for Machine Learning*, 2024.

Fengmiao Bian, Jian-Feng Cai, and Rui Zhang. A preconditioned Riemannian gradient descent algorithm for low-rank matrix recovery. *SIAM Journal on Matrix Analysis and Applications*, 45(4):2075–2103, 2024.

Vladimir Bogachev, Vladimir Aletov, Alexander Molozhavenko, Denis Bobkov, Vera Soboleva, Aibek Alanov, and Maxim Rakhuba. LoRA meets Riemannion: Muon optimizer for parametrization-independent low-rank adapters. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=WtbXgc9GVA>.

Léon Bottou. Large-scale machine learning with stochastic gradient descent. In *International Conference on Computational Statistics*, pp. 177–186. Springer, 2010.

David Carlson, Volkan Cevher, and Lawrence Carin. Stochastic spectral descent for restricted Boltzmann machines. In *International Conference on Artificial Intelligence and Statistics*, pp. 111–119, 2015a.

David Carlson, Ya-Ping Hsieh, Edo Collins, Lawrence Carin, and Volkan Cevher. Stochastic spectral descent for discrete graphical models. *IEEE Journal of Selected Topics in Signal Processing*, 10(2):296–311, 2015b.

David E Carlson, Edo Collins, Ya-Ping Hsieh, Lawrence Carin, and Volkan Cevher. Preconditioned spectral descent for deep learning. In *Advances in neural information processing systems*, volume 28, 2015c.

Yair Carmon, John C Duchi, Oliver Hinder, and Aaron Sidford. Lower bounds for finding stationary points I. *Mathematical Programming*, 184(1):71–120, 2020.

Franz Louis Cesista. Muon and a selective survey on steepest descent in Riemannian and non-Riemannian manifolds, 2025. URL <http://leloykun.github.io/ponder/steepest-descent-non-riemannian/>.

- 
- Lizhang Chen, Jonathan Li, and Qiang Liu. Muon optimizes under spectral norm constraints. *arXiv preprint arXiv:2506.15054*, 2025.
- Petros Drineas, Ravi Kannan, and Michael W Mahoney. Fast Monte Carlo algorithms for matrices I: Approximating matrix multiplication. *SIAM Journal on Computing*, 36(1):132–157, 2006a.
- Petros Drineas, Ravi Kannan, and Michael W Mahoney. Fast Monte Carlo algorithms for matrices II: Computing a low-rank approximation to a matrix. *SIAM Journal on Computing*, 36(1):158–183, 2006b.
- John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for online learning and stochastic optimization. *Journal of Machine Learning Research*, 12(7), 2011.
- Sai Surya Subramanyam Duvvuri, Fnu Devvrit, Rohan Anil, Cho-Jui Hsieh, and Inderjit S. Dhillon. Combining axes preconditioners through Kronecker approximation for deep learning. In *International Conference on Learning Representations*, 2024.
- Athanasios Glentis, Jiaxiang Li, Andi Han, and Mingyi Hong. A minimalist optimizer design for LLM pretraining. *arXiv preprint arXiv:2506.16659*, 2025.
- Vineet Gupta, Tomer Koren, and Yoram Singer. Shampoo: Preconditioned stochastic tensor optimization. In *International Conference on Machine Learning*, pp. 1842–1850, 2018.
- Nathan Halko, Per-Gunnar Martinsson, and Joel A Tropp. Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM Review*, 53(2):217–288, 2011.
- Yongchang Hao, Yanshuai Cao, and Lili Mou. Flora: Low-rank adapters are secretly gradient compressors. In *International Conference on Machine Learning*, 2024.
- Chuan He, Zhaosong Lu, Defeng Sun, and Zhanwang Deng. Complexity of normalized stochastic first-order methods with momentum under heavy-tailed noise. *arXiv preprint arXiv:2506.11214*, 2025.
- Geoffrey Hinton, Nitish Srivastava, and Kevin Swersky. Neural networks for machine learning lecture 6a overview of mini-batch gradient descent. 14(8):2, 2012.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *International Conference on Learning Representations*, 1(2):3, 2022.
- K Jordan, Y Jin, V Boza, Y Jiacheng, F Cecista, L Newhouse, and J Bernstein. Muon: An optimizer for hidden layers in neural networks. URL <https://kellerjordan.github.io/posts/muon>.
- Keller Jordan, Jeremy Bernstein, Brendan Rappazzo, @fernbear.bsky.social, Boza Vlado, You Jiacheng, Franz Cesista, Braden Koszarsky, and @Grad62304977. modded-nanogpt: Speedrunning the nanogpt baseline, 2024. URL <https://github.com/KellerJordan/modded-nanogpt>.
- Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations*, 2015.
- Dmitry Kovalev. Understanding gradient orthogonalization for deep learning via non-Euclidean trust-region optimization. *arXiv preprint arXiv:2503.12645*, 2025.
- Tim Tsz-Kit Lau, Qi Long, and Weijie Su. PolarGrad: A class of matrix-gradient optimizers from a unifying preconditioning perspective. *arXiv preprint arXiv:2505.21799*, 2025.
- Yann LeCun, Bernhard Boser, John S. Denker, Donnie Henderson, Richard E. Howard, Wayne Hubbard, and Lawrence D. Jackel. Backpropagation applied to handwritten zip code recognition. *Neural Computation*, 1(4):541–551, 1989.
- Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015.

- 
- Jiaxiang Li and Mingyi Hong. A note on the convergence of Muon and further. *arXiv e-prints*, pp. arXiv-2502, 2025.
- Liming Liu, Zhenghao Xu, Zixuan Zhang, Hao Kang, Zichong Li, Chen Liang, Weizhu Chen, and Tuo Zhao. COSMOS: A hybrid adaptive optimizer for memory-efficient training of LLMs. *arXiv preprint arXiv:2502.17410*, 2025.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.
- Chao Ma, Wenbo Gong, Meyer Scetbon, and Edward Meeds. SWAN: Preprocessing SGD enables Adam-level performance on LLM training with significant memory reduction. *arXiv e-prints*, pp. arXiv-2412, 2024.
- Sadhika Malladi, Tianyu Gao, Eshaan Nichani, Alex Damian, Jason D Lee, Danqi Chen, and Sanjeev Arora. Fine-tuning language models with just forward passes. In *Advances in Neural Information Processing Systems*, volume 36, pp. 53038–53075, 2023.
- Thomas Pethick, Wanyun Xie, Kimon Antonakopoulos, Zhenyu Zhu, Antonio Silveti-Falls, and Volkan Cevher. Training deep learning models with norm-constrained LMOs. *arXiv preprint arXiv:2502.07529*, 2025.
- Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. 2019.
- Artem Riabinin, Egor Shulgin, Kaja Grutkowska, and Peter Richtárik. Gluon: Making Muon & Scion great again! (bridging theory and practice of LMO-based optimizers for LLMs). *arXiv preprint arXiv:2505.13416*, 2025.
- Frank Rosenblatt. The perceptron: A probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6):386, 1958.
- Naoki Sato, Hiroki Naganuma, and Hideaki Iiduka. Analysis of Muon’s convergence and critical batch size. *arXiv preprint arXiv:2507.01598*, 2025.
- Maria-Eleni Sfyraiki and Jun-Kun Wang. Lions and Muons: Optimization via stochastic Frank-Wolfe. *arXiv preprint arXiv:2506.04192*, 2025.
- Wei Shen, Ruichuan Huang, Minhui Huang, Cong Shen, and Jiawei Zhang. On the convergence analysis of Muon. *arXiv preprint arXiv:2505.23737*, 2025.
- Chongjie Si, Debing Zhang, and Wei Shen. AdaMuon: Adaptive Muon optimizer. *arXiv preprint arXiv:2507.11005*, 2025.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. LLaMA: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Mark Tuddenham, Adam Prügel-Bennett, and Jonathan Hare. Orthogonalising gradients to speed up neural network optimisation. *arXiv preprint arXiv:2202.07052*, 2022.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in neural information processing systems*, volume 30, 2017.
- Nikhil Vyas, Depen Morwani, Rosie Zhao, Itai Shapira, David Brandfonbrener, Lucas Janson, and Sham M. Kakade. SOAP: Improving and stabilizing Shampoo using Adam for language modeling. In *International Conference on Learning Representations*, 2025.
- Shuo Xie, Tianhao Wang, Sashank J. Reddi, Sanjiv Kumar, and Zhiyuan Li. Structured preconditioners in adaptive optimization: A unified analysis. *arXiv preprint arXiv:2503.10537*, 2025.
- Matthew D Zeiler. Adadelta: An adaptive learning rate method. *arXiv preprint arXiv:1212.5701*, 2012.

---

Jingzhao Zhang, Sai Praneeth Karimireddy, Andreas Veit, Seungyeon Kim, Sashank Reddi, Sanjiv Kumar, and Suvrit Sra. Why are adaptive methods good for attention models? In *Advances in Neural Information Processing Systems*, volume 33, pp. 15383–15393, 2020.

Jiawei Zhao, Zhenyu Zhang, Beidi Chen, Zhangyang Wang, Anima Anandkumar, and Yuandong Tian. GaLore: Memory-efficient LLM training by gradient low-rank projection. In *International Conference on Machine Learning*, volume 235, pp. 61121–61143, 2024.



## 5 Experimental details

In this part, we present the details for our experiments on pretraining GPT-2 and LLaMA models. Table 3 presents the hyperparameter configurations for the GPT-2 and LLaMA models. We employ hyperparameters aligned with those used in (Jordan et al.; Zhao et al., 2024). In addition, we set the maximum training sequence length to 65,536 tokens, the maximum validation sequence length to 262,144 tokens, and the validation dataset size to 10,485,760 tokens. The tuning budget, token budget, validation setup, hardware, and random-seed protocol are described below.

Table 3: Hyperparameter configurations for GPT-2 and LLaMA models.

| # Parameters | Hidden Dimension | Intermediate Dimension | # Heads | # Layers |
|--------------|------------------|------------------------|---------|----------|
| 60M          | 512              | 1376                   | 8       | 8        |
| 135M         | 768              | 2048                   | 12      | 12       |
| 350M         | 1024             | 2736                   | 16      | 24       |
| 1B           | 2048             | 5461                   | 24      | 32       |

**Training, validation, and tuning protocol.** All GPT-2 and LLaMA pretraining experiments are run from random initialization without using pretrained checkpoints. We use FineWeb10B, FineWeb100B, and FineWebEdu10B, and each run lasts 5000 optimizer iterations. The batch size is 2, so each iteration processes at most  $2 \times 65,536 = 131,072$  training tokens and each full run processes at most  $5000 \times 131,072 = 655,360,000$  training tokens. Validation uses a held-out split with 10,485,760 tokens and maximum sequence length 262,144, and we report validation perplexity, computed as  $\exp(\text{validation loss})$ . The learning rate is warmed up during the first 70% of training and then follows cosine annealing to 10% of the initial rate. Hyperparameters are tuned over the Cartesian product of the grids in Table 6: 9 SGDM configurations, 9 AdamW configurations, 12 Muon configurations, 24 fixed-rank low-rank Muon configurations, and one adaptive-rank low-rank Muon configuration for each reported adaptive run. We select the best end-of-budget validation perplexity across the grid. All runs use 2 NVIDIA A100 GPUs with 80GB memory each. For the main comparison tables, we report empirical variability over three runs with different random seeds.

**Additional efficiency diagnostics.** We compute the additional efficiency diagnostics directly from the training logs. For a target perplexity  $\tau$ , the time-to-target perplexity is the accumulated training time at the first validation checkpoint satisfying  $\exp(\text{val\_loss}) \leq \tau$ , using the validation trace containing step, val\_loss, and train\_time. Tokens/sec is computed as the number of processed training tokens divided by wall-clock training time; with our training budget this equals  $5000 \times 131,072$  divided by the reported computational time in the main tables or training-time logs. Optimizer overhead is measured by timing the optimizer update separately from the forward/backward computation with CUDA events. For Muon-style methods, we report the comparable matrix-orthogonalization/update component, excluding common optimizer bookkeeping shared across variants; this makes the fixed-rank and adaptive low-rank rows directly comparable with full Muon. Peak memory is reported from the final PyTorch CUDA peak-memory counters, including both peak allocated and peak reserved memory.

Across the completed adaptive-rank runs in Table 5, the final validation perplexity ranges from 18.94 to 45.31, corresponding to validation losses from 2.94 to 3.81. The wall-clock time ranges from  $2.69 \times 10^3$  to  $2.91 \times 10^5$  seconds, the time to reach perplexity at most 50 ranges from  $1.13 \times 10^3$  to  $2.33 \times 10^4$  seconds, the throughput ranges from  $2.25 \times 10^3$  to  $2.44 \times 10^5$  tokens/sec, and peak allocated CUDA memory ranges from 13.1GB to 76.3GB.

We present the algorithmic hyperparameter search grids for each method in Table 6, following a similar choice as in (Jordan et al.; Lau et al., 2025). The best numerical performance for each training method across the grids is presented in Section 4. In addition, we apply a warm-up schedule for the learning rate during the first 70% of the training iterations, followed by a cosine annealing schedule that reduces the learning rate to 10% of its initial value.

Table 4: Representative efficiency summary across the reported GPT-2 and LLaMA pretraining tables. We report one representative value for each method. Wall-clock time is summarized from Tables 1 and 2; tokens/sec is computed as  $5000 \times 131,072$  divided by wall-clock time. Time-to-target is read from the validation-perplexity curves, using the first checkpoint that reaches PPL 50 when available and the final checkpoint otherwise. Peak memory is copied as the raw allocated/reserved MiB pair from the final PyTorch CUDA peak-memory line of representative completed logs, and optimizer overhead reports the comparable optimizer-update component; for Muon-style methods this is the matrix-orthogonalization/update overhead.

| Model | Method     | Wall-clock (s)     | Tokens/sec         | Time to PPL $\leq 50$ (s) | Peak Mem. alloc./res. (MiB) | Optimizer overhead |
|-------|------------|--------------------|--------------------|---------------------------|-----------------------------|--------------------|
| GPT-2 | SGDM       | $2.60 \times 10^4$ | $2.52 \times 10^4$ | $2.56 \times 10^4$        | 32922 / 33974               | 5%                 |
| GPT-2 | AdamW      | $2.60 \times 10^4$ | $2.52 \times 10^4$ | $1.73 \times 10^4$        | 33084 / 34074               | 12%                |
| GPT-2 | GaLore     | $2.70 \times 10^4$ | $2.42 \times 10^4$ | $1.10 \times 10^4$        | 32828 / 33522               | 23%                |
| GPT-2 | SOAP       | $2.64 \times 10^4$ | $2.48 \times 10^4$ | $3.92 \times 10^4$        | 58676 / 65480               | 32%                |
| GPT-2 | Muon       | $2.92 \times 10^4$ | $2.24 \times 10^4$ | $3.05 \times 10^4$        | 33003 / 33774               | 42%                |
| GPT-2 | LR-Muon100 | $2.79 \times 10^4$ | $2.35 \times 10^4$ | $1.41 \times 10^4$        | 33003 / 33494               | 16%                |
| GPT-2 | LR-Muon200 | $2.85 \times 10^4$ | $2.30 \times 10^4$ | $1.18 \times 10^4$        | 33003 / 33714               | 20%                |
| GPT-2 | AdaRank    | $2.80 \times 10^4$ | $2.34 \times 10^4$ | $1.29 \times 10^4$        | 32317 / 32654               | 21%                |
| LLaMA | SGDM       | $2.63 \times 10^4$ | $2.50 \times 10^4$ | $3.07 \times 10^4$        | 22018 / 24754               | 5%                 |
| LLaMA | AdamW      | $2.62 \times 10^4$ | $2.51 \times 10^4$ | $1.76 \times 10^4$        | 22144 / 24834               | 12%                |
| LLaMA | GaLore     | $2.56 \times 10^4$ | $2.56 \times 10^4$ | $1.99 \times 10^4$        | 32828 / 33522               | 23%                |
| LLaMA | SOAP       | $2.69 \times 10^4$ | $2.44 \times 10^4$ | $2.50 \times 10^4$        | 59806 / 61914               | 32%                |
| LLaMA | Muon       | $2.90 \times 10^4$ | $2.27 \times 10^4$ | $3.19 \times 10^4$        | 22081 / 23834               | 42%                |
| LLaMA | LR-Muon100 | $2.75 \times 10^4$ | $2.38 \times 10^4$ | $1.64 \times 10^4$        | 22081 / 23774               | 15%                |
| LLaMA | LR-Muon200 | $2.84 \times 10^4$ | $2.31 \times 10^4$ | $1.63 \times 10^4$        | 22081 / 23814               | 19%                |
| LLaMA | AdaRank    | $2.79 \times 10^4$ | $2.35 \times 10^4$ | $1.64 \times 10^4$        | 22270 / 23894               | 22%                |

Table 5: Completed-run efficiency diagnostics for adaptive-rank low-rank Muon runs. The final-perplexity and wall-clock entries match the AdaRank rows in Tables 1 and 2. Time-to-target uses the first validation checkpoint reaching perplexity at most 50. Tokens/sec is computed using the full-run training-token budget  $5000 \times 131,072$  divided by the final wall-clock time. Peak memory reports allocated/reserved CUDA memory in MiB.

| Model | Scale | Dataset       | Final PPL | Wall-clock (s)     | Time to PPL $\leq 50$ (s) | Tokens/sec         | Peak Mem. (MiB) |
|-------|-------|---------------|-----------|--------------------|---------------------------|--------------------|-----------------|
| GPT-2 | 60M   | FineWeb10B    | 35.61     | $2.72 \times 10^3$ | $1.52 \times 10^3$        | $2.41 \times 10^5$ | 14040 / 15712   |
| GPT-2 | 135M  | FineWeb10B    | 28.31     | $7.55 \times 10^3$ | $3.02 \times 10^3$        | $8.68 \times 10^4$ | 32317 / 32654   |
| GPT-2 | 350M  | FineWeb10B    | 26.75     | $1.25 \times 10^4$ | $3.75 \times 10^3$        | $5.24 \times 10^4$ | 48063 / 50362   |
| GPT-2 | 1B    | FineWeb10B    | 21.45     | $2.91 \times 10^5$ | $2.33 \times 10^4$        | $2.25 \times 10^3$ | 76312 / 78940   |
| GPT-2 | 60M   | FineWeb100B   | 35.93     | $2.79 \times 10^3$ | $1.56 \times 10^3$        | $2.35 \times 10^5$ | 14040 / 15712   |
| GPT-2 | 135M  | FineWeb100B   | 29.58     | $7.68 \times 10^3$ | $2.92 \times 10^3$        | $8.53 \times 10^4$ | 32317 / 32654   |
| GPT-2 | 350M  | FineWeb100B   | 26.76     | $1.08 \times 10^4$ | $3.24 \times 10^3$        | $6.07 \times 10^4$ | 48063 / 50362   |
| GPT-2 | 1B    | FineWeb100B   | 21.02     | $2.89 \times 10^5$ | $2.17 \times 10^4$        | $2.27 \times 10^3$ | 76312 / 78940   |
| GPT-2 | 60M   | FineWebEdu10B | 32.51     | $2.73 \times 10^3$ | $1.37 \times 10^3$        | $2.40 \times 10^5$ | 14040 / 15712   |
| GPT-2 | 135M  | FineWebEdu10B | 23.20     | $7.58 \times 10^3$ | $2.27 \times 10^3$        | $8.65 \times 10^4$ | 32317 / 32654   |
| GPT-2 | 350M  | FineWebEdu10B | 21.20     | $1.26 \times 10^4$ | $3.15 \times 10^3$        | $5.20 \times 10^4$ | 48063 / 50362   |
| GPT-2 | 1B    | FineWebEdu10B | 18.94     | $2.84 \times 10^5$ | $1.99 \times 10^4$        | $2.31 \times 10^3$ | 76312 / 78940   |
| LLaMA | 60M   | FineWeb10B    | 38.15     | $2.69 \times 10^3$ | $1.61 \times 10^3$        | $2.44 \times 10^5$ | 13101 / 14548   |
| LLaMA | 135M  | FineWeb10B    | 34.96     | $5.69 \times 10^3$ | $2.28 \times 10^3$        | $1.15 \times 10^5$ | 22270 / 23894   |
| LLaMA | 350M  | FineWeb10B    | 33.80     | $1.29 \times 10^4$ | $3.87 \times 10^3$        | $5.08 \times 10^4$ | 48063 / 50362   |
| LLaMA | 1B    | FineWeb10B    | 22.98     | $2.90 \times 10^5$ | $2.32 \times 10^4$        | $2.26 \times 10^3$ | 76240 / 79012   |
| LLaMA | 60M   | FineWeb100B   | 34.90     | $2.77 \times 10^3$ | $2.49 \times 10^3$        | $2.37 \times 10^5$ | 13101 / 14548   |
| LLaMA | 135M  | FineWeb100B   | 34.56     | $5.78 \times 10^3$ | $3.47 \times 10^3$        | $1.13 \times 10^5$ | 22270 / 24014   |
| LLaMA | 350M  | FineWeb100B   | 32.89     | $1.11 \times 10^4$ | $3.66 \times 10^3$        | $5.90 \times 10^4$ | 48063 / 50362   |
| LLaMA | 1B    | FineWeb100B   | 21.89     | $2.82 \times 10^5$ | $2.12 \times 10^4$        | $2.32 \times 10^3$ | 76240 / 79012   |
| LLaMA | 60M   | FineWebEdu10B | 32.30     | $2.83 \times 10^3$ | $1.13 \times 10^3$        | $2.32 \times 10^5$ | 13101 / 14548   |
| LLaMA | 135M  | FineWebEdu10B | 29.51     | $5.68 \times 10^3$ | $1.99 \times 10^3$        | $1.15 \times 10^5$ | 22270 / 23894   |
| LLaMA | 350M  | FineWebEdu10B | 28.35     | $1.27 \times 10^4$ | $3.18 \times 10^3$        | $5.16 \times 10^4$ | 48063 / 50362   |
| LLaMA | 1B    | FineWebEdu10B | 23.45     | $2.89 \times 10^5$ | $2.02 \times 10^4$        | $2.27 \times 10^3$ | 76240 / 79012   |

Table 6: Algorithmic hyperparameters. Here, NS stands for Newton–Schulz.

| Method        | Hyperparameters             | Values                                                       |
|---------------|-----------------------------|--------------------------------------------------------------|
| SGDM          | batch size                  | 2                                                            |
|               | learning rate               | $\{10^{-4}, 10^{-3}, 10^{-2}\}$                              |
|               | weight decay                | $\{0, 10^{-3}, 10^{-2}\}$                                    |
| AdamW         | batch size                  | 2                                                            |
|               | learning rate               | $\{10^{-4}, 10^{-3}, 10^{-2}\}$                              |
|               | weight decay                | $\{0, 10^{-3}, 10^{-2}\}$                                    |
| Muon          | batch size                  | 2                                                            |
|               | learning rate               | $\{10^{-2}, 2.5 \times 10^{-2}, 5 \times 10^{-2}, 10^{-3}\}$ |
|               | # NS steps                  | $\{5, 10, 20\}$                                              |
|               | weight decay                | 0                                                            |
| Low-rank Muon | batch size                  | 2                                                            |
|               | learning rate               | $\{10^{-2}, 2.5 \times 10^{-2}, 5 \times 10^{-2}, 10^{-3}\}$ |
|               | rank                        | $\{100, 200\}$                                               |
|               | # NS steps                  | $\{5, 10, 20\}$                                              |
|               | weight decay                | 0                                                            |
| AdaRank       | batch size                  | 2                                                            |
|               | hidden-matrix learning rate | $2.5 \times 10^{-2}$                                         |
|               | rank estimator              | stable-rank proxy estimated at each optimizer step           |
|               | rank range                  | $64 \leq r_k \leq 256$                                       |
|               | rank rounding               | round upward in increments of 32                             |
|               | safety factor               | $\delta_k = 0.1(k+1)^{-0.25}$                                |
|               | weight decay                | $10^{-2}$                                                    |

## 6 Appendix: Proof of the Main Results

In this section, we provide the proofs of our main results presented in Section 3, specifically Theorems 1 to 4.

We first provide three technical lemmas, whose proofs can be found in (He et al., 2025, Lemmas 1 to 3) and are therefore omitted here. The next lemma provides an expansion for the  $\alpha$ -power of the Frobenius norm, generalizing the well-known identity  $\|U + V\|_F^2 = \|U\|_F^2 + 2\langle U, V \rangle + \|V\|_F^2$  and inequality  $\|U + V\|_F^2 \leq (1 + c)\|U\|_F^2 + (1 + 1/c)\|V\|_F^2$  for all  $U, V \in \mathbb{R}^{m \times n}$  and  $c > 0$ .

**Lemma 1.** *For all  $\alpha \in (1, 2]$ ,  $U, V \in \mathbb{R}^{m \times n}$ , and  $c > 0$ , it holds that*

$$\|U + V\|_F^\alpha \leq \|U\|_F^\alpha + \alpha\|U\|_F^{\alpha-2}\langle U, V \rangle + 2\|V\|_F^\alpha, \quad (14)$$

$$\|U + V\|_F^\alpha \leq (1 + c)\|U\|_F^\alpha + (2 + (\alpha - 1)\alpha^{-1}c^{1-\alpha})\|V\|_F^\alpha. \quad (15)$$

The next lemma provides an estimation of the partial sums of series.

**Lemma 2.** *Let  $\zeta(\cdot)$  be a convex univariate function. Then we have  $\sum_{r=a}^b \zeta(r) \leq \int_{a-1/2}^{b+1/2} \zeta(\tau) d\tau$  for any integers  $a, b$  satisfying  $[a - 1/2, b + 1/2] \subset \text{dom } \zeta$ . Consequently, one has*

$$\sum_{r=a}^b r^{-\beta} \leq \begin{cases} \ln(b + \frac{1}{2}) - \ln(a - \frac{1}{2}) & \text{if } \beta = 1, \\ \frac{(b + \frac{1}{2})^{1-\beta} - (a - \frac{1}{2})^{1-\beta}}{1-\beta} & \text{if } \beta > 0, \beta \neq 1. \end{cases} \quad (16)$$

We next provide a lemma that will be used to derive complexity bounds for our methods.

**Lemma 3.** *Let  $\beta \in (0, 1)$  and  $u \in (0, 1/e)$  be given. Then,  $v^{-\beta} \ln v \leq 2u/\beta$  holds for all  $v \geq (u^{-1} \ln(1/u))^{1/\beta}$ .*

We next establish a descent property for  $f$  along a matrix-signed direction.

**Lemma 4.** *Suppose that Assumption 1 holds. Let  $X, M \in \mathbb{R}^{m \times n}$  and  $\eta > 0$  be given, and let  $X^+ = X - \eta \text{msgn}(M)$ . Then we have*

$$f(X^+) \leq f(X) - \eta \|\nabla f(X)\|_* + 2\eta \|\nabla f(X) - M\|_* + \frac{L_* \eta^2}{2}, \quad (17)$$

where  $L_*$  is given in Assumption 1.

*Proof.* By the definition of the matrix-sign function, one has  $\|\text{msgn}(M)\| \leq 1$ , and  $\langle M, \text{msgn}(M) \rangle = \|M\|_*$ . It then follows from these and (4) with  $Y = X^+$  that

$$\begin{aligned} f(X^+) &\stackrel{(4)}{\leq} f(X) + \langle \nabla f(X), X^+ - X \rangle + \frac{L_*}{2} \|X^+ - X\|^2 \\ &= f(X) + \langle M, X^+ - X \rangle + \langle \nabla f(X) - M, X^+ - X \rangle + \frac{L_*}{2} \|X^+ - X\|^2 \\ &\leq f(X) - \eta \langle M, \text{msgn}(M) \rangle + \eta \|\nabla f(X) - M\|_* \|\text{msgn}(M)\| + \frac{L_* \eta^2}{2} \|\text{msgn}(M)\|^2 \\ &\leq f(X) - \eta \|M\|_* + \eta \|\nabla f(X) - M\|_* + \frac{L_* \eta^2}{2} \\ &\leq f(X) - \eta \|\nabla f(X)\|_* + 2\eta \|\nabla f(X) - M\|_* + \frac{L_* \eta^2}{2}, \end{aligned}$$

where the second inequality is due to  $X^+ = X - \eta \text{msgn}(M)$  and the trace Hölder inequality, the third inequality is due to  $\|\text{msgn}(M)\| \leq 1$  and  $\langle M, \text{msgn}(M) \rangle = \|M\|_*$ , and the last inequality follows from the triangular inequality. Hence, the conclusion (17) holds as desired.  $\square$

## 6.1 Proof of the Main Results in Section 3.1

In this subsection, we prove Theorem 1.

**Proof of Theorem 1.** The proof of (5) can be found in (Halko et al., 2011, Theorem 10.5). Next, we prove  $\text{msgn}(QQ^T M) = Q\text{msgn}(Q^T M)$ . Let  $Q^T M = U_Q \Sigma V^T$  be the reduced SVD of  $Q^T M$ , where  $U_Q$  and  $V$  have orthogonal columns, and  $\Sigma$  is diagonal. Denote  $U = QU_Q$ . Since  $Q$  and  $U_Q$  have orthogonal columns, we obtain that  $U$  also has orthogonal columns. Thus,  $QQ^T M = U\Sigma V^T$  represents a reduced SVD of  $QQ^T M$ . Therefore, we have  $\text{msgn}(QQ^T M) = UV^T = QU_Q V^T = Q\text{msgn}(Q^T M)$ .  $\square$

## 6.2 Proof of the Main Results in Section 3.2

In this subsection, we give proofs of Theorems 2 and 3.

**Proof of Theorem 2.** Recall from Algorithm 2 that  $M_O^k = \text{msgn}(M_Q^k)$  for all  $k \geq 0$ . Using this and Lemma 4 with  $(X, X^+, M, \eta) = (X^k, X^{k+1}, M_Q^k, \eta_k)$ , we obtain that for all  $k \geq 0$ ,

$$f(X^{k+1}) \leq f(X^k) - \eta_k \|\nabla f(X^k)\|_* + 2\eta_k \|\nabla f(X^k) - M_Q^k\|_* + \frac{L_* \eta_k^2}{2}. \quad (18)$$

Summing (18) over  $k = 0, \dots, K-1$  and using  $f(X^K) \geq f_{\text{low}}$ , we obtain that for all  $K \geq 1$ ,

$$f_{\text{low}} \leq f(X^K) \stackrel{(18)}{\leq} f(X^0) - \eta_{K-1} \sum_{k=0}^{K-1} \|\nabla f(X^k)\|_* + \sum_{k=0}^{K-1} \left( 2\eta_k \|\nabla f(X^k) - M_Q^k\|_* + \frac{L_* \eta_k^2}{2} \right), \quad (19)$$

where the last inequality is due to (18) and the fact that  $\{\eta_k\}$  is nonincreasing. Rearranging this inequality and using  $\eta_k = (k+1)^{-1/2}$  for all  $k \geq 0$ , we obtain that for all  $K \geq 3$ ,

$$\begin{aligned} \frac{1}{K} \sum_{k=0}^{K-1} \|\nabla f(X^k)\|_* &\stackrel{(19)}{\leq} \frac{f(X^0) - f_{\text{low}}}{K\eta_{K-1}} + \frac{\sum_{k=0}^{K-1} \left( 2\eta_k \|\nabla f(X^k) - M_Q^k\|_* + \frac{L_* \eta_k^2}{2} \right)}{K\eta_{K-1}} \\ &= \frac{f(X^0) - f_{\text{low}}}{K^{1/2}} + \frac{\sum_{k=0}^{K-1} \left( \frac{2\|\nabla f(X^k) - M_Q^k\|_*}{(k+1)^{1/2}} + \frac{L_*}{2(k+1)} \right)}{K^{1/2}} \\ &\leq \frac{f(X^0) - f_{\text{low}} + L_* \ln K}{K^{1/2}} + \frac{2 \sum_{k=0}^{K-1} \frac{\|\nabla f(X^k) - M_Q^k\|_*}{(k+1)^{1/2}}}{K^{1/2}}, \end{aligned}$$

where the last inequality follows from  $\sum_{k=0}^{K-1} 1/(k+1) \leq \ln(2K+1) \leq 2 \ln K$  for all  $K \geq 3$  due to (16). Hence, the conclusion (6) holds as desired.  $\square$

**Proof of Theorem 3.** Using (7), (8), the definition of  $\{(\eta_k, \delta_k)\}$ , and the same arguments as for proving (6), we obtain that for all  $K \geq 3$ ,

$$\begin{aligned} \min_{0 \leq k \leq K-1} \{\|\nabla f(X^k)\|\} &\leq \frac{(f(X^0) - f_{\text{low}} + L_*) \ln K}{K^{1/2}} + \frac{2}{K^{1/2}} \sum_{k=0}^{K-1} \frac{\delta_k}{(k+1)^{1/2}} \\ &= \frac{(f(X^0) - f_{\text{low}} + L_*) \ln K}{K^{1/2}} + \frac{2}{K^{1/2}} \sum_{k=0}^{K-1} \frac{1}{k+1} \\ &\leq \frac{(f(X^0) - f_{\text{low}} + L_* + 2) \ln K}{K^{1/2}} \stackrel{(8)}{=} \frac{U_{\text{gd}} \ln K}{K^{1/2}}, \end{aligned} \quad (20)$$

where the last inequality follows from  $\sum_{k=0}^{K-1} 1/(k+1) \leq \ln(2K+1) \leq 2 \ln K$  for all  $K \geq 3$  due to (16). In addition, by Lemma 3 with  $(\beta, u, v) = (1/2, \epsilon/(4U_{\text{gd}}), K)$ , one can see that

$$K^{-1/2} \ln K \leq \frac{\epsilon}{U_{\text{gd}}} \quad \forall K \geq \left( \frac{4U_{\text{gd}}}{\epsilon} \ln \left( \frac{4U_{\text{gd}}}{\epsilon} \right) \right)^2,$$

which along with (20) implies that Theorem 3 holds.  $\square$

### 6.3 Proof of the Main Results in Section 3.3

In this subsection, we begin by establishing some technical lemmas and use them to prove Theorems 4. For convenience, we define a sequence of potentials for Algorithm 4 as follows:

$$\mathcal{P}_k := f(X^k) + p_k \|\nabla f(X^k) - M^k\|_F^\alpha \quad \forall k \geq 0, \quad (21)$$

where the sequence  $\{(X^k, M^k)\}$  is generated by Algorithm 4, and  $\{p_k\}$  is a sequence of positive scalars that will be specified separately later. The following lemma presents a recurrence relation for the estimation error of the gradient estimators  $\{M^k\}$  generated by Algorithm 4.

**Lemma 5.** *Suppose that Assumptions 1 and 2 hold. Let  $\{(X^k, M^k)\}$  be generated by Algorithm 4 with input parameters  $\{(\eta_k, \theta_k)\}$ . Then we have for all  $k \geq 0$ ,*

$$\mathbb{E}_{\xi^{k+1}} [\|M^{k+1} - \nabla f(X^{k+1})\|_F^\alpha] \leq (1 - \theta_k) \|M^k - \nabla f(X^k)\|_F^\alpha + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha + 2\sigma^\alpha \theta_k^\alpha, \quad (22)$$

where  $L_*$  is given in Assumption 1, and  $\sigma$  and  $\alpha$  are given in Assumption 2.

*Proof.* Fix any  $k \geq 0$ . It follows from (9) that

$$\begin{aligned} & M^{k+1} - \nabla f(X^{k+1}) \\ & \stackrel{(9)}{=} (1 - \theta_k)(M^k - \nabla f(X^k) + \nabla f(X^k) - \nabla f(X^{k+1})) + \theta_k(G(X^{k+1}; \xi^{k+1}) - \nabla f(X^{k+1})). \end{aligned} \quad (23)$$

In addition, we observe from (10) that  $M_Q^k = \text{msgn}(M_Q^k)$ , which along with (11) implies that  $\|X^{k+1} - X^k\| = \eta_k \|\text{msgn}(M_Q^k)\| \leq \eta_k$ . Also, it follows from Assumption 1 and Assumption 2 that  $\mathbb{E}_{\xi^{k+1}}[G(X^{k+1}; \xi^{k+1}) - \nabla f(X^{k+1})] = 0$ ,  $\mathbb{E}_{\xi^{k+1}}[\|G(X^{k+1}; \xi^{k+1}) - \nabla f(X^{k+1})\|_F^\alpha] \leq \sigma^\alpha$ , and  $\|\nabla f(X^k) - \nabla f(X^{k+1})\|_F \leq \|\nabla f(X^k) - \nabla f(X^{k+1})\|_* \leq L_* \eta_k$ . Using these, (14), (15), and (23), we obtain that for all  $c > 0$ ,

$$\begin{aligned} & \mathbb{E}_{\xi^{k+1}} [\|M^{k+1} - \nabla f(X^{k+1})\|_F^\alpha] \\ & \stackrel{(23)}{=} \mathbb{E}_{\xi^{k+1}} [\|(1 - \theta_k)(M^k - \nabla f(X^k)) + (1 - \theta_k)(\nabla f(X^k) - \nabla f(X^{k+1})) + \theta_k(G(X^{k+1}; \xi^{k+1}) - \nabla f(X^{k+1}))\|_F^\alpha] \\ & \stackrel{(14)}{\leq} \|(1 - \theta_k)(M^k - \nabla f(X^k)) + (1 - \theta_k)(\nabla f(X^k) - \nabla f(X^{k+1}))\|_F^\alpha \\ & \quad + 2\mathbb{E}_{\xi^{k+1}} [\|\theta_k(G(X^{k+1}; \xi^{k+1}) - \nabla f(X^{k+1}))\|_F^\alpha] \\ & \stackrel{(15)}{\leq} (1 + c)(1 - \theta_k)^\alpha \|M^k - \nabla f(X^k)\|_F^\alpha + L_*^\alpha (2 + (\alpha - 1)^{\alpha-1} c^{1-\alpha}) (1 - \theta_k)^\alpha \eta_k^\alpha + 2\sigma^\alpha \theta_k^\alpha, \end{aligned} \quad (24)$$

where the first inequality follows from (14) and  $\mathbb{E}_{\xi^{k+1}}[G(X^{k+1}; \xi^{k+1})] = \nabla f(X^{k+1})$ , the second inequality is due to (15),  $\mathbb{E}_{\xi^{k+1}}[\|G(X^{k+1}; \xi^{k+1}) - \nabla f(X^{k+1})\|_F^\alpha] \leq \sigma^\alpha$ , and  $\|\nabla f(X^k) - \nabla f(X^{k+1})\|_F \leq L_* \eta_k$ .

When  $\theta_k = 1$ , (22) clearly holds. For  $\theta_k \in (0, 1)$ , letting  $c = (1 - \theta_k)^{1-\alpha} - 1$  in (24), and using the fact that  $\alpha \in (1, 2]$ , we have

$$c^{1-\alpha} = ((1 - \theta_k)^{1-\alpha} - 1)^{1-\alpha} = \left( \frac{1}{(1 - \theta_k)^{\alpha-1}} - 1 \right)^{1-\alpha} \leq \left( \frac{1}{1 - (\alpha - 1)\theta_k} - 1 \right)^{1-\alpha} \leq ((\alpha - 1)\theta_k)^{1-\alpha},$$

where the first inequality follows from  $(1 - \tau)^\beta \leq 1 - \beta\tau$  for all  $\tau \in (-\infty, 1)$  and  $\beta \in [0, 1]$ . Combining this inequality with (24), one has

$$\mathbb{E}_{\xi^{k+1}} [\|M^{k+1} - \nabla f(X^{k+1})\|_F^\alpha] \leq (1 - \theta_k) \|M^k - \nabla f(X^k)\|_F^\alpha + L_*^\alpha (2 + \theta_k^{1-\alpha}) (1 - \theta_k)^\alpha \eta_k^\alpha + 2\sigma^\alpha \theta_k^\alpha,$$

which along with  $\theta_k \in (0, 1]$  and  $\alpha \in (1, 2]$  implies that (22) holds as desired.  $\square$

The following lemma establishes a descent property for the potential sequence  $\{\mathcal{P}_k\}$  defined below.

**Lemma 6.** Suppose that Assumptions 1 and 2 hold. Let  $\{(X^k, M^k)\}$  be generated by Algorithm 4 with input parameters  $\{(\eta_k, \theta_k, \delta_k)\}$ ,  $L_*$  be given in Assumption 1,  $\sigma$  and  $\alpha$  be given in Assumption 2, and  $\varrho := \max\{m, n\}$ , and let  $\{\mathcal{P}_k\}$  be defined in (21) for  $\{(X^k, M^k)\}$  and any nonincreasing positive sequence  $\{p_k\}$ . Then, it holds that for all  $k \geq 0$ ,

$$\mathbb{E}_{\xi^{k+1}}[\mathcal{P}_{k+1}] \leq \mathcal{P}_k - \eta_k \|\nabla f(X^k)\|_* + 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + \frac{(\alpha - 1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{\alpha^{\alpha/(\alpha-1)} (\theta_k p_k)^{1/(\alpha-1)}} + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k. \quad (25)$$

*Proof.* Fix any  $k \geq 0$ . Recall from (10) that  $M_O^k = \text{msgn}(M_Q^k)$  and  $\|M^k - M_Q^k\|_* \leq \delta_k$ . Using these and (17) with  $(X^+, X, M, \eta) = (X^{k+1}, X^k, M_Q^k, \eta_k)$ , we obtain that

$$\begin{aligned} f(X^{k+1}) &\leq f(X^k) - \eta_k \|\nabla f(X^k)\|_* + 2\eta_k \|\nabla f(X^k) - M_Q^k\|_* + \frac{L_* \eta_k^2}{2} \\ &\leq f(X^k) - \eta_k \|\nabla f(X^k)\|_* + 2\eta_k \|\nabla f(X^k) - M^k\|_* + 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2}. \end{aligned} \quad (26)$$

Combining this with (21) and (22), we obtain that

$$\begin{aligned} \mathbb{E}_{\xi^{k+1}}[\mathcal{P}_{k+1}] &\stackrel{(21)}{=} \mathbb{E}_{\xi^{k+1}}[f(X^{k+1}) + p_{k+1} \|\nabla f(X^{k+1}) - M^{k+1}\|_F^\alpha] \\ &\stackrel{(22)(26)}{\leq} f(X^k) - \eta_k \|\nabla f(X^k)\|_* + 2\eta_k \|\nabla f(X^k) - M^k\|_* + (1 - \theta_k) p_{k+1} \|M^k - \nabla f(X^k)\|_F^\alpha \\ &\quad + 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_{k+1} + 2\sigma^\alpha \theta_k^\alpha p_{k+1} \\ &\leq f(X^k) - \eta_k \|\nabla f(X^k)\|_* + 2\eta_k \varrho^{1/2} \|\nabla f(X^k) - M^k\|_F + (1 - \theta_k) p_k \|M^k - \nabla f(X^k)\|_F^\alpha \\ &\quad + 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k, \end{aligned} \quad (27)$$

where the second inequality follows from  $\|U\|_* \leq \varrho^{1/2} \|U\|_F$  for all  $U \in \mathbb{R}^{m \times n}$ , and the fact that  $\{p_k\}$  is nonincreasing. In addition, letting  $\alpha' = \alpha/(\alpha - 1)$  and using the Young's inequality, we obtain that

$$\begin{aligned} 2\eta_k \varrho^{1/2} \|\nabla f(X^k) - M^k\|_F &\leq \frac{((\alpha \theta_k p_k)^{1/\alpha} \|\nabla f(X^k) - M^k\|_F)^\alpha}{\alpha} + \frac{\left(\frac{2\varrho^{1/2} \eta_k}{(\alpha \theta_k p_k)^{1/\alpha}}\right)^{\alpha'}}{\alpha'} \\ &= \theta_k p_k \|\nabla f(X^k) - M^k\|_F^\alpha + \frac{(\alpha - 1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{\alpha^{\alpha/(\alpha-1)} (\theta_k p_k)^{1/(\alpha-1)}}. \end{aligned}$$

This along with (27) implies that

$$\begin{aligned} \mathbb{E}_{\xi^{k+1}}[\mathcal{P}_{k+1}] &\leq f(X^k) + p_k \|\nabla f(X^k) - M^k\|_F^\alpha - \eta_k \|\nabla f(X^k)\|_* \\ &\quad + 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + \frac{(\alpha - 1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{\alpha^{\alpha/(\alpha-1)} (\theta_k p_k)^{1/(\alpha-1)}} + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k. \end{aligned}$$

The conclusion (25) then follows from this and (21).  $\square$

We are now ready to prove Theorem 4.

**Proof of Theorem 4.** Let  $\{(X^k, M^k)\}$  be generated by Algorithm 4 with the definition of  $\{(\eta_k, \theta_k)\}$ , and  $\{\mathcal{P}_k\}$  be defined in (21) with such  $\{(X^k, M^k)\}$  and the following  $\{p_k\}$ :

$$p_k = (k + 1)^{(\alpha^2 - 3\alpha + 2)/(3\alpha - 2)} \quad \forall k \geq 0. \quad (28)$$

Since  $\alpha \in (1, 2]$ , one can see that  $\{p_k\}$  is nonincreasing. Also, observe from the definition of  $\{(\eta_k, \theta_k)\}$  that  $\{\eta_k\} \subset (0, 1]$  and  $\{\theta_k\} \subset (0, 1]$ . Hence,  $\{(\eta_k, \theta_k, p_k)\}$  satisfies the assumptions in Lemma 6 and Algorithm 4. In addition, by (21) and (28), one has that

$$\mathbb{E}[\mathcal{P}_0] = f(X^0) + \mathbb{E}[\|G(X^0; \xi^0) - \nabla f(X^0)\|_F^\alpha] \leq f(X^0) + \sigma^\alpha, \quad (29)$$

$$\mathbb{E}[\mathcal{P}_K] \geq \mathbb{E}[f(X^K)] \geq f_{\text{low}}. \quad (30)$$

Taking expectation on both sides of (25) with respect to  $\{\xi^i\}_{i=0}^{k+1}$ , we have for all  $k \geq 0$ ,

$$\begin{aligned} \mathbb{E}[\mathcal{P}_{k+1}] &\leq \mathbb{E}[\mathcal{P}_k] - \eta_k \mathbb{E}[\|\nabla f(X^k)\|_*] + 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + \frac{(\alpha-1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{\alpha^{\alpha/(\alpha-1)} (\theta_k p_k)^{1/(\alpha-1)}} \\ &\quad + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k. \end{aligned}$$

Summing up this inequality over  $k = 0, \dots, K-1$ , and using (29) and (30), we obtain that for all  $K \geq 1$ ,

$$\begin{aligned} f_{\text{low}} &\stackrel{(30)}{\leq} \mathbb{E}[\mathcal{P}_K] \leq \mathbb{E}[\mathcal{P}_0] - \sum_{k=0}^{K-1} \eta_k \mathbb{E}[\|\nabla f(X^k)\|_*] \\ &\quad + \sum_{k=0}^{K-1} \left( 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + \frac{(\alpha-1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{\alpha^{\alpha/(\alpha-1)} (\theta_k p_k)^{1/(\alpha-1)}} + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k \right) \end{aligned} \quad (31)$$

$$\begin{aligned} &\stackrel{(29)}{\leq} f(X^0) + \sigma^\alpha - \eta_{K-1} \sum_{k=0}^{K-1} \mathbb{E}[\|\nabla f(X^k)\|_*] \\ &\quad + \sum_{k=0}^{K-1} \left( 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + \frac{(\alpha-1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{\alpha^{\alpha/(\alpha-1)} (\theta_k p_k)^{1/(\alpha-1)}} + 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k \right), \end{aligned} \quad (32)$$

where the last inequality follows from (29) and the fact that  $\{\eta_k\}$  is nonincreasing. Rearranging the terms in (32), and using the definition of  $\{(\eta_k, \theta_k, \delta_k)\}$ , (12), and (28), we obtain that for all  $K \geq 3$ ,

$$\begin{aligned} \frac{1}{K} \sum_{k=0}^{K-1} \mathbb{E}[\|\nabla f(X^k)\|_*] &\stackrel{(32)}{\leq} \frac{f(X^0) - f_{\text{low}} + \sigma^\alpha}{K \eta_{K-1}} \\ &\quad + \frac{\sum_{k=0}^{K-1} \left( 2\eta_k \delta_k + \frac{L_* \eta_k^2}{2} + \frac{(\alpha-1)(2\varrho^{1/2} \eta_k)^{\alpha/(\alpha-1)}}{(\theta_k p_k)^{1/(\alpha-1)}} \right)}{K \eta_{K-1}} + \frac{\sum_{k=0}^{K-1} \left( 3L_*^\alpha \theta_k^{1-\alpha} \eta_k^\alpha p_k + 2\sigma^\alpha \theta_k^\alpha p_k \right)}{K \eta_{K-1}} \\ &\stackrel{(28)}{=} \frac{f(X^0) - f_{\text{low}} + \sigma^\alpha}{K^{(\alpha-1)/(3\alpha-2)}} + \frac{\sum_{k=0}^{K-1} \frac{L_*}{2(k+1)^{2(2\alpha-1)/(3\alpha-2)}}}{K^{(\alpha-1)/(3\alpha-2)}} + \frac{\sum_{k=0}^{K-1} \frac{2+(\alpha-1)(2\varrho^{1/2}/\alpha)^{\alpha/(\alpha-1)} + 3L_*^\alpha + 2\sigma^\alpha}{k+1}}{K^{(\alpha-1)/(3\alpha-2)}} \\ &\stackrel{(12)}{\leq} \frac{U_{\text{mn}} \ln K}{K^{(\alpha-1)/(3\alpha-2)}}, \end{aligned}$$

where the second inequality follows from  $\sum_{k=0}^{K-1} 1/(k+1) \leq 2 \ln K$  due to (16) and  $K \geq 3$ , and  $\sum_{k=0}^{K-1} 1/(k+1)^{2(2\alpha-1)/(3\alpha-2)} \leq (3\alpha-2)2^{\alpha/(3\alpha-2)}/\alpha < 4$  due to (16) and  $(3\alpha-2)/\alpha \in (1, 2]$ . Recall that  $\iota_K$  is uniformly selected from  $\{0, \dots, K-1\}$ . It then follows from this and the above inequality that for all  $K \geq 3$ ,

$$\mathbb{E}[\|\nabla f(X^{\iota_K})\|_*] \leq \frac{U_{\text{mn}} \ln K}{K^{(\alpha-1)/(3\alpha-2)}}. \quad (33)$$

In addition, by Lemma 3 with  $(\beta, u, v) = ((\alpha-1)/(3\alpha-2), (\alpha-1)\epsilon/(2(3\alpha-2)U_{\text{mn}}), K)$ , one can see that

$$K^{-(\alpha-1)/(3\alpha-2)} \ln K \leq \frac{\epsilon}{U_{\text{mn}}} \quad \forall K \geq \left( \frac{2(3\alpha-2)U_{\text{mn}}}{(\alpha-1)\epsilon} \ln \left( \frac{2(3\alpha-2)U_{\text{mn}}}{(\alpha-1)\epsilon} \right) \right)^{(3\alpha-2)/(\alpha-1)},$$

which together with (33) implies that Theorem 4 holds.  $\square$